

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



High Level Programming

The IO Library

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

License Information

This work is licensed under the license



Attribution-NonCommercial-NoDerivatives 4.0 International

This license requires that reusers give credit to the creator. It allows reusers to copy and distribute the material in any medium or format in unadapted form and for noncommercial purposes only.

① **BY:** Credit must be given to you, the creator.

② **NC:** Only noncommercial use of your work is permitted.

Noncommercial means not primarily intended for or directed towards commercial advantage or monetary compensation.

③ **ND:** No derivatives or adaptations of your work are permitted.

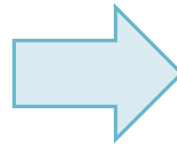
To view a copy of the license, visit:

<https://creativecommons.org/licenses/by-nc-nd/4.0/?ref=chooser-v1>

Premises

❖ Where are we?

- u01-courseIntroduction
- u02-review
- u03-cppBasics
- u04-cppLibrary**
- u05-multithreading
- u06-synchronization
- u07-advancedIO
- u08-IPC



- u04s01e
- u04s03e
- u04s04e
- u04s07e
- u04s08e
- u04s01-IOLibrary.pdf
- u04s02-sequentialContainers.pdf
- u04s03-associativeContainers.pdf
- u04s04-genericAlgorithms.pdf
- u04s05-dynamicMemory.pdf
- u04s06-copyControl.pdf
- u04s07-templates.pdf
- u04s08-exercise.pdf

```
for c++: #include <fstream>
#include <string>
```

```
int main() {
    std::ifstream f;
    f.open("filename");
    if(!f.is_open()) {
        return -1; // return error
    } else {
        std::string a, b, c;
        while(f >> a >> b) {
            c = a + b;
            // do something with c
        }
    }
    f.close();
    return 0;
}
```

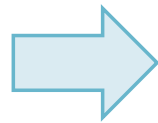
```
#include <fcntl.h>
#include <unistd.h>

int main() {
    char buffer[100];
    int fd = open("filename", O_RDONLY);
    if (fd == -1) {
        // handle error
        return -1;
    } else {
        ssize_t n = read(fd, buffer, sizeof(buffer) - 1);
        if (n == -1) {
            // handle error
            close(fd);
            return -1;
        } else {
            buffer[n] = '\0'; // null-terminate the string
            // do something with buffer
        }
        close(fd);
    }
    return 0;
}
```

Condition states

- ❖ When we perform IO, an error can occur
 - Some errors are recoverable, others are not
 - Once an error has occurred, subsequent IO operations will fail
 - The code should always check for errors in IO streams
 - The easiest way to check the state of a stream object is to use a condition

```
cin >> s;
```



```
if (cin >> s) { ... Success ... }  
while (cin >> s) { ... Success ... }
```

IO Library

- ❖ The header **fstream** defines the types used to read and write named files
 - These types provide the same operations as those we have used on the objects **cin** and **cout**

```
#include <fstream> (header)
```

```
std::ifstream file("example.txt")
```

```
std::ifstream inputFile("output.txt");
```

```
std::ofstream outputFile("output.txt");
```

- Thanks to inheritance, we can use the standard IO operators on them
 - **<<**, **>>**, and **getline**

Type	Meaning
ifstream	To read from a given file.
ofstream	To write to given file.
fstream	To read or write a given file.

IO Library

- ❖ To support languages that use wide characters, the library defines a set of objects manipulating `wchart_t`
 - These objects' names start with
 - `Ifstream` → `wifstream`, etc.
 - `cin` → `wcin`, `cout` → `wcout`, `cerr` → `wcerr`
- ❖ Notice that IO objects cannot be assigned

```
ofstream fo1, fo2;  
...  
fo2 = fo1;           // Error: cannot assign streams
```


Condition states

- ❖ The IO classes also define functions and flags that we can check to understand the status of the stream
 - If **s** is a stream we can interrogate its status

Type	Meaning
s.eof()	True if the stream hits EOF. <small>End of file reached.</small>
s.fail()	True if the IO operation failed. <small>error reading from the file</small>
s.bad()	True if the stream is corrupted. <small>Unrecoverable error occurred.</small>
s.good()	True if the stream is in a valid state.
s.clear()	Reset all condition values to a valid state. Return void.
etc.	

Attention:
while (!s.eof()) { ... }
is buggy !!!

IO Library

- ❖ We can also use specific members to manage the file associated with the stream

Type	Meaning
<code>fstream fs;</code>	Creates an unbound file stream. <small>fs is created but not associated with any file</small>
<code>fstream fs(f);</code>	Creates an fstream fs and open file f. The object f must be a string or a C-like pointer to a string. <small>fs is associated with the f file.</small>
<code>fstream fs(f,mode);</code>	Like the previous one but open the file in the given mode. <small>read or write mode: <code>std::ios::in</code> <code>std::ios::out</code></small>
<code>fs.open(f)</code>	Open the file f and bound it to fs.
<code>fs.open(f,mode)</code>	Like the previous one, but open the file in the given mode. Please see the next page.
<code>fs.close()</code>	Close the file to which fs is bound. Return void.
<code>fs.is_open()</code>	Return a bool to indicate whether the file associated with fs is open.

fs: Stream
f: File name

fstream or
ifstream or
ofstream

IO Library

- ❖ Each stream has an associated file **mode**
 - The mode can be changed when the file is opened
 - In output mode, the previous content is lost (if we do not append)

Type	Meaning
in	Open for input.
out	Open for output.
app	Append: Seek to the end before every write. All writes are at the end of the file. <small>When opening a file with this mode, the output will be appended to the end of the file. This means that the file pointer will be positioned at the end of the file before every write operation, ensuring that data is added to the existing content.</small>
ate	Seek to the end immediately after open. Then, it is possible to move around (seek). Combine it with other types.
trunc	Truncate the file. In output, the default it truncate the file (rewrite it) even if trunc is not specified. <small>When opening a file in output mode (out or app), specifying this mode will truncate (clear) the contents of the file before any write operations. If trunc is not specified and the file already exists, its contents will be truncated by default.</small>
binary	Perform IO operation in binary mode. Combine it wioth other types.

Examples

Opening a file in
different mode

```
// Out and trunc are implicit
ofstream out("myfile");

// Trunc is implicit
ofstream out("myfile", ofstream::out);

ofstream out("myfile", ofstream::out | ofstream::trunc);

ofstream out("myfile", ofstream::out | ofstream::app);

// Out is implicit
ofstream out("myfile", ofstream::app);
```

Examples

```
#include <iostream>
#include <fstream>
#include <string>

typedef struct my_s {
    string title;
    int x, y;
} my_t;

my_t test;
test.title = ""; test.x = test.y = 0;
ofstream file;

file.open ("example.txt", ofstream::out);
file << test.title;
file << "\t" << test.x;
file << "\t" << test.y << std::endl;

file.close();
```

Write some file
content

endl: writes a newline and flushes the buffer
flush: flushes the buffer
ends: writes a NULL and flushes the buffer

Examples

```
#include <iostream>
#include <fstream>

using std::cout; using std::cerr;
using std::endl; using std::string;
using std::ifstream;

int main() {
    string filename("input.txt");
    int number;

    ifstream infile(filename);
    if (!infile.is_open()) {
        cerr << "Error: " << filename << endl;
        return EXIT_FAILURE;
    }
    ...
    infile.close();

    return EXIT_SUCCESS;
}
```

Opening a file in
reading mode

Main reading cycle
... see ahead

Examples

All writes are flushed immediately
`setbuf(stdout,0);` → `cout << unitbuf;`
Opposite operation:
`cout << nunitbuf;`

```
int number;  
  
cout << unitbuf;  
  
while (infile >> number) {  
    cout << number << " ";  
}  
cout << endl;
```

Read integer
values

```
while (infile.get(c)) {  
    cout << c << " ";  
}  
cout << endl;
```

Read single
characters

Examples

Does not read
(and print)
newlines

```
string s;  
while (getline(infile,s)) {  
    cout << s << "; ";  
}  
cout << endl;
```

Read entire file
lines


```
string s;  
while (infile >> s) {  
    cout << s << "; ";  
}  
cout << endl;
```

Read single
strings

Exercise: IO in C, UNIX, and C++

❖ An IO library comparison

- Write a segment of C code that writes and then reads a sequence of N integer values
- Use the ASCII and binary forms
 - C library
 - UNIX library
 - C++ library
- Compare file size and runtimes



```
#include <iostream>
#include <fstream>
#include <cstdlib>

#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#include <unistd.h>
#include <fcntl.h>

using std::cout;
using std::endl;
using std::fstream;
using std::ofstream;
using std::ifstream
```

Solution

Time library

Computing times

```
#include <chrono>
```

```
long long microseconds;
```

```
auto start = std::chrono::high_resolution_clock::now();
```

```
write();
```

```
OR
```

```
read();
```

Either write or read the file

```
auto elapsed = std::chrono::high_resolution_clock::  
    now() - start;
```

```
microseconds = std::chrono::duration_cast<std::chrono::  
    microseconds>(elapsed).count();
```

```
cout << " C    Write Time: " << microseconds << endl;
```

Solution

The C library is buffered

```
setbuf(fp, 0);  
fp = fopen ("...", "w");  
for (i=0; i<n; i++) {  
    val = rand();  
    fprintf (fp, "%d\n", val);  
}  
fclose (fp);
```

C Buffered

Force a buffer size
equal to zero

```
fp = fopen ("...", "r");  
while (fscanf (fp, "%d", &val) != EOF) {  
    i++;  
}  
fclose (fp);
```

Solution

UNIX

```
fd = open ("...", O_WRONLY | O_CREAT | ...),  
for (i=0; i<n; i++) {  
    val = rand();  
    write (fd, &val, sizeof (int));  
}  
close (fd);
```

```
fd = open ("...", O_RDONLY);  
while (read (fd, &val, sizeof (int)) != 0) {  
    i++;  
}  
close (fd);
```

Solution

Trade-off Time-Memory

UNIX +

```
v = (int *) malloc (n * sizeof (int));  
fd = open ("...", O_WRONLY | O_CREAT | ...);  
for (i=0; i<n; i++) {  
    v[i] = rand();  
}  
write (fd, v, n*sizeof (int));  
close (fd);
```

```
fd = open ("...", O_RDONLY);  
read (fd, v, n*sizeof (int));  
close (fd);
```

Solution

C++

```
s.open ("...", ofstream::out);  
for (i=0; i<n; i++) {  
    val = rand();  
    s << val << endl;  
}  
s.close();
```

Force an fflush
Unbuffered

```
s.open ("...", ifstream::in);  
while (s >> val) {  
    i++;  
}  
s.close();
```


Solution

C++ Binary

```
s.open ("...",  
        ofstream::out | ofstream::binary);  
for (i=0; i<n; i++) {  
    val = rand();  
    s.write ((char *) &val, sizeof (int));  
}  
s.close();
```

```
s.open ("...", ifstream::in | ifstream::binary);  
while (s.read((char *) &val, sizeof (int))) {  
    i++;  
}  
s.close();
```

Solution

❖ File Size [MBytes]

Variable
(due to the size of
random integers)

$$n = 2 \cdot 10^6$$

Library	ASCII	Binary
C, UNIX, C++	20.965	8.00

❖ CPU Time [seconds]

Fixed

Library	Write	Read
C (ASCII, buffered)	0.172	0.125
UNIX (binary)	1.513	0.411
C++ (ASCII, unbuffered)	1.762	0.102

Solution

❖ File Size [MBytes]

Variable
(due to the size of
random integers)

$$n = 2 \cdot 10^6$$

Library	ASCII	Binary
C, UNIX, C++	20.965	8.00

❖ CPU Time [seconds]

Fixed

Library	Write	Read
C (ASCII, buffered)	0.172	0.125
C (ASCII, unbuffered)	1.806	0.139
UNIX (binary)	1.513	0.411
UNIX (binary, array)	0.028	0.001
C++ (ASCII, unbuffered)	1.762	0.102
C++ (binary)	0.046	0.022