

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
```

```
{
    int freq[MAXPAROLA]; /* vettore di conteggi
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```

It's like a set of guidelines that helps software developers write programs that can run on many different types of Unix-like systems without needing major changes.

POSIX stands for Portable Operating System Interface, a set of standards that define the interface between a unix-like os and application software. The goal is to ensure compatibility and interoperability between unix like systems.

POSIX standards are primarily associated with Unix and Unix-like operating systems such as Linux, macOS, and various flavors of BSD

For example, if a programmer wants to create a new process, read from a file, or manage memory, POSIX tells them exactly how to do it in a standardized way. This makes it easier for software to be portable, meaning it can be used on different systems without a lot of extra work.

# Multi-Threading

## Multi-Threading POSIX

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

# Pthreads

Pthreads (POSIX) is the standard library for handling threads in unix-like os. It defines a set of functions and data structures for creating, controlling, and synchronizing threads.

## ❖ POSIX threads or Pthreads

### ➤ Is the standard UNIX library for threads

- The version POSIX 1003.1c was born in 1995
- Revised in version IEEE POSIX 1003.1-2017

### ➤ Defined for the C language, but available for other languages (e.g., FORTRAN)

## ❖ Using Pthreads

### ➤ A thread is a **function** that is executed in concurrency with the main thread

A process with multiple threads is a set of independent that share the process resources

# Pthreads

- ❖ The Pthreads library allows
  - Creating and manipulating threads
  - Synchronizing threads
  - Protection of resources shared by threads
  - Thread scheduling
  - Destroying thread
- ❖ It defines more than 60 functions
  - All functions have the prefix **pthread\_**
    - pthread\_equal, pthread\_self, pthread\_create, pthread\_exit, pthread\_join, pthread\_cancel, pthread\_detach

## Library linkage

- ❖ The Pthread system calls are defined in `pthread.h`
  - Insert in all `*.c` files
    - `#include <pthread.h>`
  - Link programs with the pthread library
    - `gcc -Wall -g -o <exeName> <file.c> -lpthread`

# Thread Identifier

- ❖ A thread is uniquely identified
  - By a type identifier **pthread\_t**
    - Similar to the PID of a process (**pid\_t**)
  - The type **pthread\_t** is opaque
    - Its definition is implementation dependent
    - Can be used only by functions specifically defined in Pthreads
    - It is not possible compare directly two identifiers or print their values
  - It has meaning only within the process where the thread is executed
    - Remember that the PID is global within the system

## System call pthread\_equal

```
int pthread_equal (
    pthread_t tid1,
    pthread_t tid2
);
```

The pthread\_equal function is a useful tool for comparing two thread identifiers (pthread\_t values) to determine if they represent the same thread.

- ❖ Compares two thread identifiers
- ❖ Arguments
  - Two thread identifiers
- ❖ Returned values
  - Nonzero if the two threads are equal
  - Zero otherwise

## System call `pthread_create`

```
pthread_t pthread_self (  
    void  
);
```

`pthread_self` returns the thread identifier of the calling thread. It's often used by a thread to identify itself within the program.

- ❖ Returns the thread identifier of the calling thread
  - It can be used by a thread (with `pthread_equal`) to self-identify

Self-identification can be important to properly access the data of a specific thread



# System call pthread\_create

```
int pthread_create (
    pthread_t *tid,
    const pthread_attr_t *attr,
    void *(*startRoutine) (void *),
    void *arg
);
```

pthread\_create is used to create a new thread. It takes several arguments including a pointer to a pthread\_t variable where the thread identifier of the newly created thread will be stored, attributes for the thread (or NULL for default attributes), a function pointer to the starting routine of the new thread, and an argument to pass to the starting routine.

Function Pointer: The starting routine is specified as a function pointer of type void \*(\*startRoutine)(void \*). This means it's a pointer to a function that takes a single void \* argument and returns a void \* result.

Return value:  
0, on success  
error code, on failure

## ❖ Arguments

➤ Identifier of the generated thread (**tid**)

➤ Thread attributes (**attr**)

- NULL is the default attribute

➤ C function executed by the thread (**startRoutine**)

➤ Argument passed to the start routine (**arg**)

- NULL if no argument

Starting routine refers to the function that will be executed by the newly created thread when it starts running. This function, referred to as the starting routine, contains the code that the new thread will execute. It's the entry point for the new thread's execution. Whatever functionality you want the new thread to perform should be implemented within this function. Any data that needs to be passed can be done through void \*arg argument of the pthread

A **single** argument

```
void *thread_function(void *arg) {
    int *thread_id_ptr = (int *)arg;
    int thread_id = *thread_id_ptr;
    printf("Hello from thread %d!\n",
        thread_id);
    // Thread's logic...
    return NULL;
}
```

```
int main() {
    pthread_t thread_id;
    int thread_arg = 1; // Example
    argument
    pthread_create(&thread_id,
        NULL, thread_function,
        &thread_arg);
    // Additional logic...
    pthread_join(thread_id, NULL); //
    Wait for the thread to finish
    return 0;
}
```



## System call `pthread_exit`

- ❖ A whole process (with all its threads) terminates if
  - Its thread calls **exit** (or **\_exit** or **\_Exit**)
  - The main thread execute **return**
  - The main thread receives a signal whose action is to terminate
- ❖ A single thread can terminate (without affecting the other process threads)
  - Executing **return** from its start function
  - Executing **pthread\_exit**
  - Receiving a cancellation request performed by another thread using **pthread\_cancel**

# System call pthread\_exit

```
void pthread_exit (
    void *valuePtr
);
```

When a thread calls pthread\_exit, it immediately terminates, and the control returns to the thread that created it or to the main thread if it's the main thread itself.

Thread Termination: pthread\_exit is used by a thread to terminate itself. When called, the thread stops executing and returns a termination status.

valuePtr: This argument allows the thread to return a termination status or value. It's a pointer to the value that the thread wants to pass to the thread that calls pthread\_join when it waits for this thread to terminate.

- ❖ It allows a thread to terminate returning a termination status
- ❖ Arguments
  - The **ValuePtr** value is kept by the kernel until a thread calls **pthread\_join**
  - This value is available to the thread that calls **pthread\_join**

# Example

Thread creation  
of 1 thread  
without  
parameters

```
void *tF () {  
    ...  
    pthread_exit (NULL);  
}
```

Attributes

Arguments

```
pthread_t tid;  
int rc;  
rc = pthread_create (&tid, NULL, tF, NULL);  
if (rc) {  
    // Error ...  
    exit (-1);  
}  
...  
pthread_exit (NULL);  
// exit (0);  
// return (0); (in main)
```

Terminates only  
the main thread

This program prints,  
Main thread is running  
Thread is running.

In practice, the order in which threads are scheduled to run can depend on various factors, including the operating system's scheduling algorithm, system load, and other concurrent activities happening on the system.

So, while the main thread is typically scheduled to run first, it's not guaranteed, and the actual order of execution may vary.

This code creates a simple multithreaded program with one thread. The main thread creates a new thread using `pthread_create`, and both threads print a message indicating their execution. Finally, both threads terminate using `pthread_exit(NULL)`.

Terminates the  
process  
(all its threads)

# Example

Creation of N threads with 1 argument

A thread can be executed when t is changed

```
void *tF (void *par) {  
    int *tidP, tid;  
    ...  
    tidP = (int *) par;  
    tid = *tidP;  
    ...  
    pthread_exit (NULL)  
}
```

The content is being modified by the main thread

```
pthread_t th[NUM_THREADS];  
int rc, t;
```

```
for (t=0; t<NUM_THREADS; t++) {  
    rc = pthread_create (&th[t], NULL, tF,  
        (void *) &t);  
    if (rc) {...}  
}  
pthread_exit(NULL);
```

## ERROR

&t is the address of a variable, the main thread changes its content in concurrency with the created threads that read its value

# Example

Creation of N  
threads with 1  
argument

Cast of a value  
`void *`  $\leftrightarrow$  **`long int`**

```
void *tF (void *par) {  
    long int tid;  
    ...  
    tid = (long int) par;  
    ...  
    pthread_exit(NULL);  
}
```

```
pthread_t th[NUM_THREADS];  
int rc; long int t;
```

```
for (t=0; t<NUM_THREADS; t++) {  
    rc = pthread_create (&th[t], NULL, tF,  
        (void *) t);  
    if (rc) { ... }  
}  
pthread_exit (NULL);
```

Tricky:

We pass a `long int` as it were an address, because `pthread_create` requires an address as its last argument

# Example

Creation of N  
threads with 1  
struct

```
struct tS {  
    int tid;  
    char str[N];  
};
```

```
void *tF (void *par) {  
    struct tS *tD;  
    int tid; char str[L];
```

```
    tD = (struct tS *) par;  
    tid = tD->tid; strcpy (str, tD->str);  
    ...
```

Cast to a vector  
of structs

```
pthread_t t[NUM_THREADS];  
struct tS v[NUM_THREADS];  
...  
for (t=0; t<NUM_THREADS; t++) {  
    v[t].tid = t;  
    strcpy (v[t].str, str);  
    rc = pthread_create (&t[t], NULL, tF, (void *) &v[t]);  
    ...  
}  
...
```

Address of a struct

## System call `pthread_join`

### ❖ At its creation a thread can be declared

#### ➤ Joinable

- Another thread may "wait" (**`pthread_join`**) for its termination, and collect its exit status
- The termination status of the thread is retained until another thread performs a **`pthread_join`** for that thread

#### ➤ Detached

- No thread can explicitly wait for its termination (not joinable)
- The termination status of the thread is immediately released



## System call pthread\_join

```
int pthread_join (  
    pthread_t tid,  
    void **valuePtr  
);
```

Return value:  
0, on success  
error code, on failure

### ❖ Arguments

- Identifier (tid) of the waited-for thread
- The void pointer **ValuePtr** will be the value returned by thread **tid**
  - Returned by **pthread\_exit** or by **return**
  - **PTHREAD\_CANCELED** if the thread was deleted
  - Can be set to NULL if you are not interested in the return value

# Example

Returns the exit status  
(**tid** in this example)

```
void *tF (void *par) {  
    long int tid;  
    ...  
    tid = (long int) par;  
    ...  
    pthread_exit ((void *) tid);  
}
```

```
void *status;  
long int s;  
...  
/* Wait for threads */  
for (t=0; t<NUM_THREADS; t++) {  
    rc = pthread_join (th[t], &status);  
    s = (long int) status;  
    if (rc) { ... }  
}  
...
```

**th[t]** collects the **tids**

Wait for each thread,  
and collects its exit  
status

## System call `pthread_cancel`

```
int pthread_cancel (  
    pthread_t tid  
);
```

Return value:  
0, on success  
error code, on failure

- ❖ Terminates the target thread
- ❖ The thread calling **pthread\_cancel** does not wait for termination of the target thread (it continues immediately after the calling)
- ❖ Arguments
  - Target thread (tid) identifier

## System call `pthread_detach`

```
int pthread_detach (  
    pthread_t tid  
);
```

Return value:  
0, on success  
error code, on failure

- ❖ Declares thread **tid** as detached
  - The status information will not be kept by the kernel at the termination of the thread
  - No thread can join with that thread
    - Calls to **pthread\_join** should fail
- ❖ Arguments
  - Thread (tid) identifier

# Example

Create a thread and  
then make it detached

```
pthread_t tid;  
int rc;  
void *status;
```

```
rc = pthread_create (&tid, NULL, PrintHello, NULL);  
if (rc) { ... }
```

```
pthread_detach (tid);
```

Detach a thread

```
rc = pthread_join (tid, &status);  
if (rc) {  
    // Error  
    exit (-1);  
}
```

Error if try to join

```
pthread_exit (NULL);
```

## Example

Create a detached thread using the attribute field

```
pthread_attr_t attr;  
void *status;
```

```
pthread_attr_init (&attr);  
pthread_attr_setdetachstate (&attr,  
    PTHREAD_CREATE_DETACHED);  
//PTHREAD_CREATE_JOINABLE);
```

Creates a detached thread

```
rc = pthread_create (&t[t], &attr, tF, NULL);  
if (rc) {...}
```

```
pthread_attr_destroy (&attr);
```

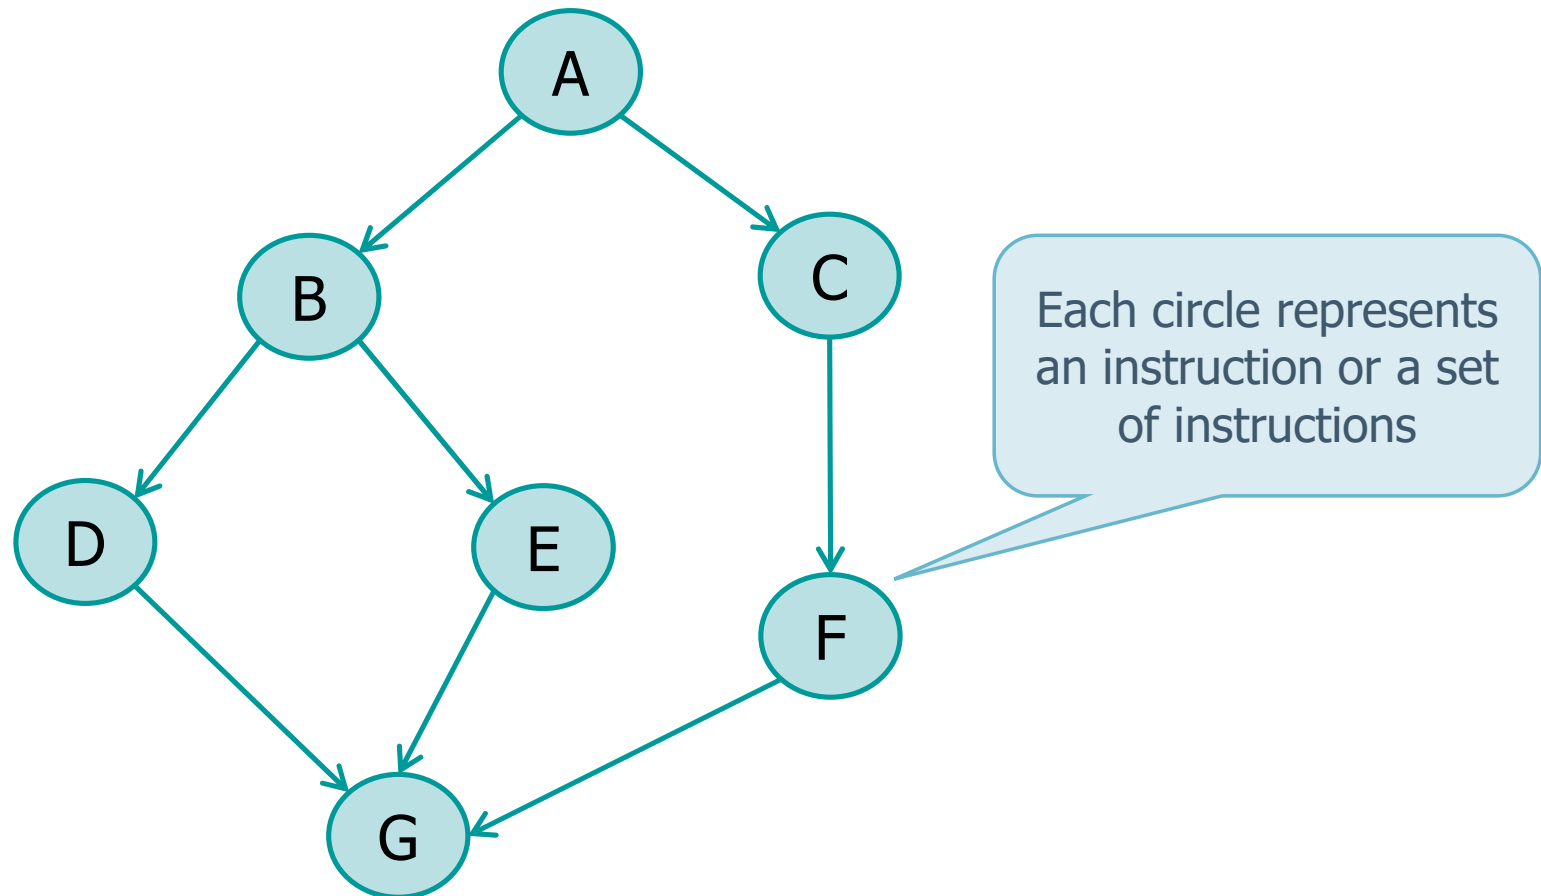
Destroys the attribute object

```
rc = pthread_join (thread[t], &status);  
if (rc) {  
    // Error  
    exit (-1);  
}
```

Error if try to join

## Exercise

- ❖ Implement, using threads, the following precedence graph using threads






# Solution

```
void waitRandomTime (int max){  
    sleep ((int)(rand() % max) + 1);  
}
```


```
int main (void) {  
    pthread_t th_cf, th_e;  
    void *retval;
```

```
    srand (getpid());  
    waitRandomTime (10);  
    printf ("A\n");
```



# Solution


```
waitRandomTime (10);  
pthread_create (&th_cf, NULL, CF, NULL);  
waitRandomTime (10);  
printf ("B\n");  
waitRandomTime (10);  
pthread_create (&th_e, NULL, E, NULL);  
waitRandomTime (10);  
printf ("D\n");  
pthread_join (th_e, &retval);  
pthread_join (th_cf, &retval);  
waitRandomTime (10);  
printf ("G\n");  
  
return 0;  
}
```



# Solution

```
static void *CF () {  
    waitRandomTime (10);  
    printf ("C\n");  
    waitRandomTime (10);  
    printf ("F\n");  
    return ((void *) 1); // Return code  
}
```

```
static void *E () {  
    waitRandomTime (10);  
    printf ("E\n");  
    return ((void *) 2); // Return code  
}
```




## Exercise

- ❖ Given a text file, with an undefined number of characters, passed as an argument of the command line
- ❖ Implement a concurrent program using three threads (**T<sub>1</sub>**, **T<sub>2</sub>**, **T<sub>3</sub>**) that process the file content in pipeline
  - **T<sub>1</sub>**: Read from file the next character
  - **T<sub>2</sub>**: Transforms the character read by **T<sub>1</sub>** in uppercase
  - **T<sub>3</sub>**: Displays the character produced by **T<sub>2</sub>** on standard output

# Solution

❖ Implement, using threads, this precedence graph



T are created and destroyed at each iteration

For now, the only synchronization strategy is to use `pthread_join()`

# Solution

```
static void *GET (void *arg) {
    char *c = (char *) arg;
    *c = fgetc (fg);
    return NULL;
}

static void *UPD (void *arg) {
    char *c = (char *) arg;
    *c = toupper (*c);
    return NULL;
}

static void *PRINT (void *arg) {
    char *c = (char *) arg;
    putchar (*c);
    return NULL;
}
```

# Solution

```
FILE *fg;

int main (int argc, char ** argv) {
    char next, this, last;
    int retC;
    pthread_t tGet, tUpd, tPrint;
    void *retV;

    if ((fg = fopen(argv[1], "r")) == NULL) {
        perror ("Error fopen\n");
        exit (0);
    }
    this = ' ';
    last = ' ';
    next = ' ';
```



## Solution

The first two characters  
can be managed  
separately

```
while (next != EOF) {
    retC = pthread_create (&tGet, NULL, GET, &next);
    if (retC != 0) fprintf (stderr, ...);
    retC = pthread_create (&tUpd, NULL, UPD, &this);
    if (retC != 0) fprintf (stderr, ...);
    retC = pthread_create (&tPrint, NULL, PRINT, &last);
    if (retcode != 0) fprintf (stderr, ...);
    retC = pthread_join (tGet, &retV);
    if (retC != 0) fprintf (stderr, ...);
    retC = pthread_join (tUpd, &retV);
    if (retC != 0) fprintf (stderr, ...);
    retC = pthread_join (tPrint, &retV);
    if (retC != 0) fprintf (stderr, ...);
    last = this;
    this = next;
}
```

## Solution

Management of the last two characters (queue)

```
// Last two chars processing
retC = pthread_create(&tUpd, NULL, UPD, &this);
if (retC!=0) fprintf (stderr, ...);
retC = pthread_create(&tPrint, NULL, PRINT, &last);
if (retC != 0) fprintf (stderr, ...);
retC = pthread_join (tUpd, &retV);
if (retC != 0) fprintf (stderr, ...);
retC = pthread_join (tPrint, &retV);
if (retC != 0) fprintf (stderr, ...);
retC = pthread_create(&tPrint, NULL, PRINT, &this);
if (retC != 0) fprintf (stderr, ...);
return 0;
}
```