

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
```

```
{
```

```
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
    for(i=0; i<MAXPAROLA; i++)
        freq[i]=0;
```

```
    if(argc != 2)
```

```
    {
        fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
        exit(1);
    }
```

```
    f = fopen(argv[1], "r");
    if(f==NULL)
```

```
    {
        fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
        exit(1);
    }
```

```
    while( fgets( riga, MAXRIGA, f ) != NULL )
```



Synchronization

Barriers

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

License Information

This work is licensed under the license



Attribution-NonCommercial-NoDerivatives 4.0 International

This license requires that reusers give credit to the creator. It allows reusers to copy and distribute the material in any medium or format in unadapted form and for noncommercial purposes only.

① **BY:** Credit must be given to you, the creator.

② **NC:** Only noncommercial use of your work is permitted.

Noncommercial means not primarily intended for or directed towards commercial advantage or monetary compensation.

③ **ND:** No derivatives or adaptations of your work are permitted.


To view a copy of the license, visit:

<https://creativecommons.org/licenses/by-nc-nd/4.0/?ref=chooser-v1>

Barriers

- ❖ Barriers can be used to coordinate multiple threads working in parallel
 - A barrier allows each thread to wait until all cooperating threads have reached the same point, and then continue executing from there

o ensure that all threads have completed a certain task before moving on to the next task. This is useful in scenarios where threads need to synchronize their progress.



Barriers

❖ Barriers generalize the `pthread_join` function

➤ The function **`pthread_join`** acts as a barrier to allow **one** thread to wait until another thread completed their processing

➤ Barriers allow **an arbitrary number** of threads to wait until all of the threads have completed their processing

➤ The threads don't have to exit, as they can continue working after all threads have reached the barrier

Generalization of

`pthread_join`:

`pthread_join`: This function

allows one thread to wait

for another thread to finish.

Think of it as waiting for

one friend to catch up

before you can continue.

Barriers: Instead of waiting

for just one friend, you

wait for all your friends.

Barriers allow multiple

threads to wait until all of

them have reached the

barrier point.

Advantages of Barriers:

Threads don't have to stop

or exit after reaching the

barrier. They can continue

working on new tasks

once all threads have

synchronized at the barrier.

Trivial solution

Semaphores: Think of semaphores as signals or flags that threads use to communicate with each other.

❖ A possible trivial solution

- Use one semaphore for each thread T_i
- Implement one for the extra process B (for the barrier) using one more semaphore

It uses too many semaphores

Each thread has its own semaphore.
There is an additional semaphore for the barrier process

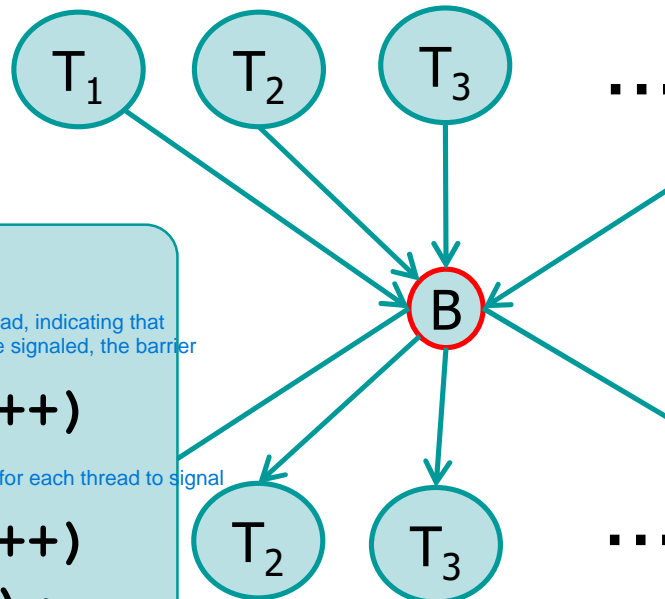
The barrier process waits n times on its semaphores and then wakes-up all threads T_i

B

The barrier process waits for a signal from each thread, indicating that they have reached the barrier. Once all threads have signaled, the barrier process signals all threads to continue.

```
for (i=0; i<n; i++)
    wait (semB); // Wait for each thread to signal
for (i=0; i<n; i++)
    signal (sem[i]);
```

// Signal all threads to continue



T_i

Barrier Process:
The barrier process waits for a signal from each thread, indicating that they have reached the barrier. Once all threads have signaled, the barrier process signals all threads to continue.

```
signal (semB);
wait (sem[i]);
```

Signal the barrier process

Waiting ...

Wait for the barrier process to signal back

Each T_i wakes-up the barrier process and wait on its own semaphore

POSIX versus C++

- ❖ In the POSIX standard barriers are implemented using the primitives

These functions are used to create and use barriers in POSIX-compliant systems (like many Unix-based systems).

- `pthread_barrier_init`, `pthread_barrier_wait`

- ❖ In C++ standard barriers are implemented by the class template `std::barrier`

- Unlike **`std::latch`**, barriers are reusable

Difference: Unlike `std::latch`, which is a one-time use synchronization point, `std::barrier` can be reused. Once all threads have synchronized, the barrier can be used again

- Once a group of arriving threads are unblocked, the barrier can be reused

- ❖ We will use the POSIX implementation

- See the documentation for the C++ version

POSIX barriers

For more details see the reference documentation

These are implemented in the pthread.h library in POSIX-compliant systems.

❖ In the POSIX standard barriers are implemented in pthread.h

Type	Meaning
<p>Purpose: Initializes a barrier. Parameters: Takes a barrier object, attributes (usually NULL), and a count of the number of threads that must reach the barrier before any can proceed. Usage: Can be initialized multiple times, but be mindful of the current value of the counter.</p> <pre>int pthread_barrier_init (...);</pre>	<p>The count argument specifies the number of threads that must reach the barrier before all of the threads will be allowed to continue. The same barrier can be initialized more than once (but pay attention to the current value of the counter).</p>
<p>Once the specified number of threads have called this function, all threads can proceed.</p> <pre>int pthread_barrier_wait (...);</pre>	<p>Indicate that a thread is done with its work and it is ready to wait for all the other threads to catch up.</p>
<pre>int pthread_barrier_destroy (...);</pre> <p>De-initializes the barrier and frees any resources allocated for it.</p>	<p>De-initialize a barrier. Any resource allocated for the barrier will be freed.</p>

Example

```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>
#define N 4 Number of threads
#define C 5 some other value c
```

```
...
```

```
pthread_barrier_t bar;
```

```
...
```

```
pthread_barrier_init (&bar, NULL, N);
```

Initializes the barrier bar with N threads. The second parameter is for attributes, which is NULL here.

```
for (i=0; i<N; i++) {
```

```
    v[i] = i;
```

th is an array of thread identifiers. Each element in the array represents a thread that is created and managed by the program.

```
    pthread_create (&th[i], NULL, f, (void *) &v[i]);
```

```
}
```

pthread_t is a data type used to uniquely identify a thread in POSIX threads (pthreads).
th is an array of pthread_t with N elements, where N is the number of threads defined by #define N 4.

```
for (i=0; i<N; i++) {
```

```
    pthread_join (th[i], NULL);
```

```
}
```

```
pthread_barrier_destroy(&bar);
```

Main program calling the threads

Define the barrier

Init the barrier with a NULL attribute and a counter equal to N

Destroy the barrier

A loop creates N threads. Each thread runs the function f. v[i] is passed as an argument to the thread function f.

(void *) &v[i] is the argument passed to the thread function f. Here, v[i] is cast to a void pointer.

Waits for all threads to finish

Why are we passing void pointer: `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void *), void *arg);` The fourth parameter, void *arg, is a pointer to the argument that will be passed to the thread function. Casting to void *:

`(void *) &v[i]` casts the address of v[i] to a void pointer.

This is necessary because the thread function f is defined to take a void * argument.

Example

f is a function that takes a void pointer and returns a void pointer.

```
void *f (void *par) {  
    int *np, n;
```

np is a pointer to an integer.
n is an integer.

```
    np = (int *) par;
```

```
    n = *np;    Casts the void pointer par to an integer pointer and dereferences it to get the value.
```

```
    fprintf (stdout, "T%d-In\n", n);
```

```
    pthread_barrier_wait(&bar);
```

```
    fprintf (stdout, "    T%d-Out\n", n);
```

```
    pthread_exit (NULL);
```

```
}
```

Use the barrier to
synchronize N threads
once

The thread function f is defined to take a void * argument and return a void *.

This is required by the pthread_create function, which expects the thread function to have this signature.

Parameter Handling:

void *par is the argument passed to the thread function.

Inside the function, par is cast to the appropriate type:

This cast converts the void * back to an int *, allowing the function to access the integer value.

Dereferencing:

n = *np; dereferences the integer pointer to get the actual integer value.

Prints a message before and after the barrier wait.

pthread_barrier_wait(&bar) makes the thread wait at the barrier until all threads reach this point.

Example

```
void *f (void *par) {
    int i, *np, n;    Adds a loop counter i.
```

```
    np = (int *) par;
```

```
    n = *np;
```

```
    for (i=0; i<C; i++) {
```

```
        fprintf (stdout, "T%d-In%d\n", n, i);
```

```
        pthread_barrier_wait(&bar);
```

```
        fprintf (stdout, "    T%d-Out%d\n", n, i);
```

```
    }
```

```
    pthread_exit (NULL);
```

```
}
```

Use the barrier to
synchronize N threads
C times

A loop runs C times.
Each iteration prints a message before and after the barrier wait.
pthread_barrier_wait(&bar) makes the thread wait at the barrier in each iteration.

Summary
POSIX Barriers: Use pthread_barrier_init, pthread_barrier_wait, and pthread_barrier_destroy to manage barriers.
Example Code: Demonstrates how to initialize, use, and destroy a barrier.
Function f: Shows how threads can synchronize using barriers, with one example synchronizing once and another synchronizing multiple times.
I hope this helps! Let me know if you have any more questions or need further clarification.

The barrier does not have
to be re-initialized

Barrier Initialization: pthread_barrier_init(&bar, NULL, N) initializes the barrier for N threads.

Barrier Wait: pthread_barrier_wait(&bar) makes threads wait until all N threads reach this point.

Barrier Destruction: pthread_barrier_destroy(&bar) cleans up the barrier.

Thread Function: The function f demonstrates how threads can synchronize using barriers, with one example synchronizing once and another synchronizing multiple times without reinitializing the barrier.

Conclusions

- ❖ Barriers are used to coordinate multiple threads working in parallel
 - You want all threads to wait until everyone has arrived at a certain point
 - A simple semaphore would do the exact opposite, i.e., each thread would keep running and the last one will go to sleep

Exercise


- ❖ Suppose barrier constructs do not exist
 - Re-implement them using only **one** semaphore **and** one mutex

```
pthread_barrier_t b;  
...  
pthread_barrier_init (&b, NULL, N_THREAD);  
for (i=0; i<N_THREAD; i++) {  
    err = pthread_create (&tid[i], NULL, thr_fn, NULL);  
}  
...
```

```
void *thr_fn () {  
    ...  
    pthread_barrier_wait (&b);  
}
```

Threads
(**acyclic** behavior)

Synchronization point
among all threads



Solution

Initializazion

```
typedef struct barrier_s {  
    sem_t sem;  
    pthread_mutex_t mutex;  
    int count;  
} barrier_t;
```

Barrier structure
sem to enqueue threads
mutex to protect counter
counter to count threads up

Init barrier

```
barrier_d = (barrier_t *) malloc (1 * sizeof(barrier_t));  
sem_init (&barrier_d->sem, 0, 0);  
pthread_mutex_init (&barrier_d->mutex, NULL);  
barrier_d->count = 0;
```


Main

```
for (i=0; i<N_THREAD; i++) {  
    err = pthread_create (&tid[i], NULL, thr_fn, NULL);  
}
```

Run threads

Solution 1

Threads
(**acyclic** behavior)



```
void *thr_fn () {  
    ...  
    pthread_mutex_lock (&barrier->mutex);  
    barrier->count++;  
    if (barrier->count == N_THREAD) {  
        for (j=0; j<N_THREAD; j++) {  
            sem_post (&barrier_d->sem);  
        }  
    }  
    pthread_mutex_unlock (&barrier->mutex);  
    sem_wait (&barrier_d->sem);  
  
    pthread_exit ();  
}
```

Protect counter

Last thread
awakes all

Un-protect counter

Waiting point for
all threads

Solution 2

Solution with turnstile

```
void *thr_fn () {  
    ...  
    pthread_mutex_lock (&barrier->mutex);  
    barrier->count++;  
    if (barrier->count == N_THREAD) {  
        sem_post (&barrier_d->sem);  
    }  
    pthread_mutex_unlock (&barrier->mutex);  
    sem_wait (&barrier_d->sem);  
    sem_post (&barrier_d->sem);  
  
    pthread_exit ();  
}
```

Protect counter

Un-protect counter

Turnstile

One **extra** sem_post is
done (pay attention to
cycling threads)

Exercise

- ❖ Re-implement the following (**cyclic**) piece of code using only one semaphore and one mutex


```
pthread_barrier_t b;  
...  
pthread_barrier_init (&b, NULL, N_THREAD);  
for (i=0; i<N_THREAD; i++) {  
    err = pthread_create (&tid[i], NULL, thr_fn, NULL);  
}  
...
```

Main

```
void *thr_fn () {  
    while (1)  
        ...  
        pthread_barrier_wait (&b);  
}  
}
```

Threads
(**cyclic** behavior)

Synchronization point
among all threads



Buggy Solution

Buggy attempt

```
void *thr_fn () {  
    while (1) {  
        ..  
        pthread_mutex_lock (&barrier->mutex);  
        barrier->count++;  
        if (barrier->count == N_THREAD) {  
            for (j=0; j<N_THREAD; j++) {  
                sem_post (&barrier_d->sem);  
            }  
        }  
        pthread_mutex_unlock (&barrier->mutex);  
        sem_wait (&barrier_d->sem);  
    }  
    pthread_exit ();  
}
```

Last threads
awakes all

Waiting point for
all threads

A fast threads can
cycle more than
once !

Solution

Initializazion

```
typedef struct barrier_s {  
    sem_t sem1, sem2;  
    pthread_mutex_t mutex;  
    int count;  
} barrier_t;
```

Barrier structure
2 sems to enqueue threads
mutex to protect counter
counter to count threads up

Init barrier

```
barrier_d = (barrier_t *) malloc (1 * sizeof(barrier_t));  
sem_init (&barrier_d->sem1, 0, 0);  
sem_init (&barrier_d->sem2, 0, 0);  
pthread_mutex_init (&barrier_d->mutex, NULL);  
barrier_d->count = 0;
```

Main

```
for (i=0; i<N_THREAD; i++) {  
    err = pthread_create (&tid[i], NULL, thr_fn, NULL);  
}
```

Run threads

Threads
(cyclic behavior)

Barrier #1

```
...  
pthread_mutex_lock (&barrier->mutex);  
barrier->count++;  
if (barrier->count == N_THREAD) {  
    for (j=0; j<N_THREAD; j++) sem_post (&barrier_d->sem1);  
}  
pthread_mutex_unlock (&barrier->mutex);  
sem_wait (&barrier_d->sem1);
```

```
pthread_mutex_lock (&barrier->mutex);  
barrier->count--;  
if (barrier->count == 0) {  
    for (j=0; j<N_THREAD; j++) sem_post (&barrier_d->sem2);  
}  
pthread_mutex_unlock (&barrier->mutex);  
sem_wait (&barrier_d->sem2);  
...
```

Barrier #2

Exercise


Concurrent Bubble-sort

- ❖ Write a version of the exchange (bubblesort) sorting algorithm) as follows
 - A static array include n integer values
 - We want to sort it using n identical threads
 - Each thread is in charge of sorting two adjacent elements
 - Thread 0 sort elements 0 and 1
 - Thread 1 sort elements 1 and 2
 - ...
 - Thread $n-1$ sort elements $n-1$ and n

Exercise

➤ Each thread

- Compare the two elements it deals with, and exchange them if they are not in the correct order
- Once their work is finished, all the threads wait for each-other, and if
 - All the elements are correctly ordered, the program terminates
 - Otherwise, all threads are run again to make a new series of exchanges




Exercise

➤ Each thread

- Compare the two elements it deals with, and exchange them if they are not in the correct order
- Once their work is finished, all the threads wait for each-other, and if
 - All the elements are correctly ordered, the program terminates
 - Otherwise, all threads are run again to make a new series of exchanges

As the order in which all swaps are performed is not defined (inner iteration) the number of necessary outer iterations is upper bounded by n

```
for (i=0; i<n-1; i++)  
    for (j=0; j<n-i-1; j++)  
        if (v[j] > v[j+1])  
            swap (v, i, j+1);
```



Solution 1

Solution in C with
semaphores (no barriers)

```
#include <stdio.h>
```

```
typedef enum {false, true} boolean;
```

```
int num_threads;
```

```
int vet_size;
```

```
int *vet;
```

```
boolean sorted = false;
```

```
boolean all_ok = false;
```

```
sem_t semMaster;
```

```
sem_t *semSlave;
```

```
pthread_mutex_t *me;
```

```
static int max_random (int);
```



```
void *master (void *);
```

```
void *slave (void *);
```

Boolean type
(implicit in C++)

Global variables:
1 semaphore for the master thread
1 semaphore for each slave thread
1 mutex for each element of the vector

Prototypes



Solution 1

Main (Extract) Part 1

```
int main (int argc, char **argv) {
```

```
    ... Definitions ...
```

```
    vet_size = atoi (argv[1]);  
    num_threads = vet_size - 1;
```


```
    ... Allocations ...
```

```
    for (i=0; i<vet_size; i++) {  
        vet[i] = max_random (1000);  
    }
```

```
    for (i=0; i<vet_size; i++) {  
        pthread_mutex_init (&me[i], NULL);  
    }
```

Fill the vector with random numbers

Create a mutex for each element of the vector



Solution 1

Main (Estract)
Part 2

MT starts

```
sem_init (&semMaster, 0, num_threads);  
pthread_create (&thMaster, NULL, master, &num_threads);
```

Creates 1 master thread

```
for (i=0; i<num_threads; i++) {
```

```
    id[i] = i;
```

STs wait

```
    sem_init (&semSlave[i], 0, 0);
```

```
    pthread_create (&thSlave[i], NULL, slave, &id[i]);
```

```
}
```

Creates num_threads
slave threads

```
for (i=0; i<num_threads; i++) {
```


```
    pthread_join (thSlave[i], NULL);
```

```
}
```

Wait all

```
pthread_join (thMaster, NULL);
```

... Free memory and semaphores ...



Solution 1

```

void *master (void *arg) {
    int *ntp, nt, i;
    ntp = (int *) arg;
    nt = *ntp;
    while (!sorted) {
        for (i=0; i<nt; i++)
            sem_wait (&semMaster);

        if (all_ok) {
            sorted = true;
        } else {
            all_ok = true;
        }



        for (i=0; i<nt; i++)
            sem_post (&semSlave[i]);
    }
    pthread_exit (0);
}

```

Wait for slave threads

If a worker performs a swap, it sets all_ok to false and here we set it back to true. If no worker performs a swap, all_ok remains true, we set sorted to true, and slaves will stop at the next iteration

Wake up slave threads



Solution 1

```

void *slave (void *arg) {
    int i = *((int *) arg);
    while (1) {
        sem_wait (&semSlave[i]);
        if (sorted) break;
        pthread_mutex_lock(&me[i]);
        pthread_mutex_lock(&me[i+1]);
        if (vet[i] > vet[i + 1]) {
            swap (vet[i], vet[i + 1]);
            all_ok = false;
        }
        pthread_mutex_unlock(&me[i+1]);
        pthread_mutex_unlock(&me[i]);
        sem_post (&semMaster);
    }
    pthread_exit (0);
}

```



Wait the
master
thread

Get 2 elements
to compare and sort

Sort them: If we do, set
all_ok to false (cycle again)

Release elements

Wake up master thread




Solution 2

Solution in C with barriers


```
#include <stdio.h>
#include <sys/timeb.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <semaphore.h>

#define N 10

int count, vet[N];
int sorted = 0;
int all_ok = 1;
sem_t me[N];
sem_t mutex, barrier1, barrier2;
```



Instead of using one semaphore for each slave, why do not use barriers?



Solution 2

read_array read or generate the array


```
int main (int argc, char * argv[]) {  
    ...  
    count = 0;  
    sem_init (&mutex, 0, 1);  
    sem_init (&barrier1, 0, 0);  
    sem_init (&barrier2, 0, 0);  
  
    for (i=0; i<N; i++)  
        sem_init (&me[i], 0, 1);  
  
    for (i=0; i<N-1; i++) {  
        id[i] = i;  
        pthread_create (&th[i], NULL, sorter, &id[i]);  
    }  
  
    pthread_exit (0);  
}
```

Create a mutex to protect the counter, and 2 barriers based on semaphores

Create a semaphore for each element of the vector

No joins (threads are detached)

Create N threads



Solution 2

```
static void *sorter (void *arg) {  
    int *a = (int *) arg;  
    int i, j, tmp;  
  
    i = *a;  
  
    pthread_detach (pthread_self ());  
  
    while (!sorted) {  
        sem_wait (&me[i]);  
        sem_wait (&me[i+1]);  
        if (vet[i] > vet[i+1]) {  
            swap (vet[i], vet[i + 1]);  
            all_ok = 0;  
        }  
        sem_post (&me[i + 1]);  
        sem_post (&me[i]);  
    }  
}
```


We do not need the master any more !!!

Acquires the 2 elements it has to manage

Sort them

all_ok remains 1 if no thread makes an exchange

Release the array elements



Solution 2

```
sem_wait (&mutex);  
count++;  
if (count == N-1) {  
    for (j=0; j<N-1; j++)  
        sem_post (&barrier1);  
}  
sem_post (&mutex);  
  
sem_wait (&barrier1);
```


Barrier #1

Before the iteration, you need to synchronize all the threads

The last thread to arrive unblock all threads

All the other threads wait on a barrier

The mutex protect the counter



Solution 2

Barrier #2

```

sem_wait (&mutex);
count--;
if (count == 0) {
    printf ("all_ok %d\n", all_ok);
    for (j=0; j<N; j++)
        printf ("%d ", vet[j]);
    printf ("\n");
    if (all_ok)
        sorted = 1;
    all_ok = 1;
    for (j=0; j<N-1; j++)
        sem_post (&barrier2);
}
sem_post (&mutex);
sem_wait (&barrier2);
}
return 0;
}

```

Restart (if necessary)


Block everything

Only one barrier is not enough, because the last thread wake up all the threads, and a fast thread can iterate more times

For this reason a second barrier is used

The last thread to arrive unblock all

All the other threads wait on a barrier



Solution 3

- ❖ Can we use `pthread_barrier_wait`?
 - Yes, and one barrier should suffice for synchronization purposes
 - Unfortunately, we must also check if the array is sorted after the barrier
 - Thus, we need a second barrier anyway
 - In the second barrier, the last thread arriving checks the sorting and displays the array
 - It is convenient to implement it with a counter, a mutex, and a semaphore

Solution 3

Solution in C with one
library barriers

```
pthread_barrier_wait (&barrier1);
```

Barrier #1

```
sem_wait (&mutex);
```

```
count++;
```

```
if (count == N-1) {
```

```
    printf ("all_ok %d\n", all_ok);
```

```
    for (j=0; j<N; j++)
```

```
        printf ("%d ", vet[j]);
```

```
    fprintf (stdout, "\n");
```

```
    if (all_ok)
```

```
        sorted = 1;
```

```
    all_ok = 1;
```

```
    for (j=0; j<N-1; j++) {
```

```
        sem_post (&barrier2);
```

```
    }
```

```
    count = 0;
```

```
}
```

```
sem_post (&mutex);
```

```
sem_wait (&barrier2);
```

Barrier #2

As Solution 2, but ...

... we must reset the
counter back to zero

