

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



# Synchronization

## Task Programming in C++

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

## License Information

This work is licensed under the license



### Attribution-NonCommercial-NoDerivatives 4.0 International

This license requires that reusers give credit to the creator. It allows reusers to copy and distribute the material in any medium or format in unadapted form and for noncommercial purposes only.

① **BY:** Credit must be given to you, the creator.

② **NC:** Only noncommercial use of your work is permitted.

Noncommercial means not primarily intended for or directed towards commercial advantage or monetary compensation.

③ **ND:** No derivatives or adaptations of your work are permitted.

To view a copy of the license, visit:

<https://creativecommons.org/licenses/by-nc-nd/4.0/?ref=chooser-v1>

# Introduction

## ❖ Multi-threading in C++ has **two** main **limitations**

1. The number of software threads may be higher than the number of hardware threads
  - **Over-subscription** occurs every time the number of software threads ready to start is higher than the number of hardware threads available in the system
  - Over-subscription implies system overhead and some performance penalty
  - With threads, the global load must be managed manually



- To solve this problem, C++11 introduced **task-based** parallel programming

# Introduction

## ❖ Multi-threading in C++ has two main limitations

### 2. Threads (the `std::thread` library) do not offer any direct way to return a value to the caller

- In POSIX
  - A simple strategy is return a value with `pthread_exit`
  - More general strategies must be user-implemented (through global or local objects)
- In native C++ only the general strategy is available (you must manipulate objects explicitly)



- To solve this problem, C++11 introduced **futures** and **promises**

## Task processing

- ❖ A task is an entity that runs asynchronously producing output data that will become available (and useful) at a later time
  - The operating system associate a thread to a task in an automatic way
  - Balancing tasks is automatic, through work-stealing features
  - Tasks have the possibility of handling return values

## Task processing

### ❖ In C++

- Thread-based parallel programming relies on **std::thread** objects

```
std::thread t(thread_function);
```

- Task-based parallel programming relies on **std::async** objects

```
auto fut = std::async(thread_function);
```



# Task processing

```
#include <future>
```

Parameters for the thread function

```
future<T> async(policy, function, args...);
```

<T> is the type of the future

Asynchronous policy

"Thread" function

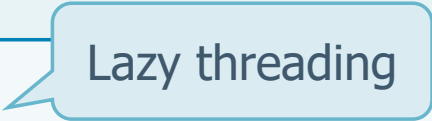
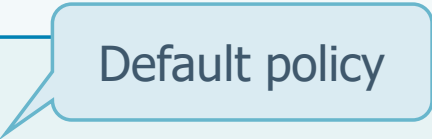
## ❖ Function `async` (namespace `std`)

- Is an alternative to `std::thread` to execute functions in parallel
- Has an extra parameter, i.e., the policy
- Returns a future of type `T`

For now, ignore it

## Task processing

- ❖ The user may decide the running policy
  - There are three different types of policies

Policy	Description
<code>launch::async</code>	Asynchronous launch, i.e., a new thread is generated to run the new function.
<code>launch::deferred</code> 	The call to the new function is deferred. The OS may never run it. The new function will be run when we <b>wait</b> it or <b>get</b> its future.
<code>launch::async   launch::deferred</code> 	The policy to run the new thread is selected by the system accordingly to the availability of concurrency in the system. It is implementation dependent.



# Examples

## Running async tasks

```
auto f1 = std::async(std::launch::async, my_f, 10);  
// Thread function my_f is run in a new thread
```

```
auto f2 = std::async(  
    std::launch::deferred, my_f, 20);  
// Thread function my_f is not run until we get  
// its results or wait for it
```

For now, we do not know what a future is

```
auto f3 = std::async(  
    std::launch::async | std::launch::deferred,  
    my_f, 30);  
// The system decides when running my_f.  
// Possibly, it never runs my_f.
```

```
f2.wait()  
// Invoke deferred function f2 (i.e., run it)
```

Force task f2 to be associated and run within a thread is

# Example

Running policy

```
auto f = std::async (my_f);
```

If it is async or deferred;  
it may never run

```
if (f.wait_for(0s)==std::future_status::deferred) {
```

```
    f.wait();
```

It is deferred: Use wait to  
force the execution.  
It is async, it is already  
running.

Wait for 0 seconds, i.e., do  
not wait, check the status

```
} else {
```

```
    while (f.wait_for(100ms) !=  
           std::future_status::ready) {
```

Check status every  
100 msecons

```
        ... do something ...
```

If it is not ready, do  
something in parallel

```
    }
```

```
    ...
```

```
}
```

Here the future f is ready

To wait for 0s or 100ms  
use std::literals

## Futures

- ❖ An **async** object will eventually hold the return value of the thread function in a future
  - A future is an object that can represent a value generated by some provider
  - Function **<future>::get** applied to a valid future
    - Blocks the thread until the object is ready
    - Returns the object (return with "return") once it is ready

<T> is the type  
of the future

```
future<T> async(policy, function, args...);
```

# Example

```
#include <future>
```

```
...
```

```
bool is_prime (int n) {  
    if (num <= 1) return false;  
    if (num <= 3) return true;  
    ...  
    return false;  
}
```

Check if num is  
prime

```
int main () {  
    std::future<bool> fut = std::async(  
        std::launch::async, is_prime, 117);  
  
    // ... do other work ...
```

Run a new function  
in a thread

```
bool ret = fut.get();  
cout << ret;
```

Wait for function is\_prime to  
return and make the Boolean  
value available

```
return 0;
```

```
}
```

# Example

```
#include <future>
#include <iostream>

...
auto fut = std::async (
    std::launch::async,
    [] () {
        std::vector<int> v;
        for (int i=0; i<100; i++)
            v.push_back(i);
        return v;
    }
);
...
auto ret = fut.get();
for (auto e: ret)
    std::cout << e << std::endl;
```

Run a new function  
thread

Work with **lambda**  
expressions

Wait for the future to be ready  
and **get** the return value

## Shared futures

- ❖ In C++ there are two types of future
  - Unique future, i.e., `std::future<T>`
    - There is only one instance referring to the event
  - Shared future, i.e., `std::shared_future<T>`
    - A `shared_future` object behaves like a `future` object, except that it can be copied
    - Multiple instances may refer to the same event
    - All instances will become ready at the same time and can be retrieved
    - May be used to signal multiple threads simultaneously, similarly to `std::condition_variable::notify_all`



# Example

## Unique future

```
int sum(int a, int b) {  
    std::this_thread::sleep_for(std::chrono::seconds(2));  
    return a + b;  
}
```

```
int main() {  
    std::future<int> fut =  
        std::async(std::launch::async, sum, 10, 20);  
    ...
```

Wait and then get the future

```
    int result1 = fut.get();  
    std::cout << "Result: " << result1 << endl;
```

Result: 30

```
    int result2 = fut.get();  
    std::cout << "Result: " << result2 << endl;
```

```
    return 0;
```

```
}
```

terminate called after throwing an instance of 'std::future\_error'  
what(): std::future\_error: No associated state  
Aborted (core dumped)

# Example

## Shared future

```
int sum(int a, int b) {  
    std::this_thread::sleep_for(std::chrono::seconds(2));  
    return a + b;  
}  
  
int main() {  
    std::shared_future<int> fut =  
        std::async(std::launch::async, sum, 10, 20);  
    ...  
  
    int result1 = fut.get();  
    std::cout << "Result: " << result1 << endl;  
  
    int result2 = fut.get();  
    std::cout << "Result: " << result2 << endl;  
  
    return 0;  
}
```

Wait and then get the future

Result: 30

Result: 30

## Promises

- ❖ The most common situation where you encounter a **future** is with a call to an **async**
  - An **async** returns a future
  - A future represents a value that you do not yet have but will have eventually
- ❖ At the lowest level, a **future** comes from an associated **promise**
  - A promise is an object that can store a value to be retrieved by a future object
    - A promise is an object that you will eventually set
  - When the value is set, it will be made available through its corresponding future

# Promises

❖ Think of promise and future as creating a single-use channel for data

➤ A promise

- Creates the channel
- Writes data in the channel

Promise  
Name

```
std::promise<type> pn;  
pn.set_value(...);
```

➤ A future

- Connects to the other end of the channel
- Waits and reads the data once it has been written

Future  
Name

```
auto fn = pn.get_future();  
fn.get();
```

# Promises

## ❖ The principal steps are

### ➤ The main thread

- Defines a promise
- Associate a future to the promise

```
std::promise<type> pn;  
auto fn = pn.get_future();
```

### ➤ The working thread

- Receive the promise
- Executes the function and fulfills the promise

```
pn.set_value(...);
```

### ➤ The main thread retrieves the result

```
fn.get();
```

# Promises

## ❖ Further considerations

- If we destroy the promise without setting a value, an exception is stored in the object
  - Function **get** will return
- The object associated to a promise is usually stored in the heap as it cannot be stored
  - In the setter of the promise, as the setter can die
  - In the getter of the future, as we futures can be shares among several getters





# Example

Promise and future.  
One thread (lambda)

Thread function

```
int f(int x) { return x + 1; }
```

Define the promise  
This is the producer-write end

```
std::promise<int> promise;  
auto future = promise.get_future();
```

Define the future from  
the promise  
This is the consumer-  
read end

```
// Launch f asynchronously  
std::thread thread([&promise] (int x) {  
    promise.set_value(f(x));  
}, 5  
);
```

Get the promise in  
the capture list

Set the value (to be  
communicated) into the promise

```
std::cout << future.get() << std::endl;
```

Get the value

## Example

Promise and future.  
Two threads (lambdas)

Define the promise and the  
future from the promise

```
auto promise = std::promise<std::string>();  
auto future = promise.get_future();
```

```
auto producer = std::thread([&] {  
    promise.set_value("Hello World");  
});
```

Run the thread  
setting the promise

```
auto consumer = std::thread([&] {  
    std::cout << future.get();  
});
```

Run the thread  
getting the future

```
producer.join();  
consumer.join();
```

## Example

Promise and future.  
One thread (function)

```
#include <future>
using namespace std;

void factorial (const int &N, promise<int>& pr) {
    int res = 1;
    for (int i=N; i> 1; i--)
        res *=i;
    pr.set_value(res);
}

int main () {
    promise<int> p;
    future<int> f = p.get_future();
    thread t = thread(factorial, 4, ref(p));
    // here we have the data
    int x = f.get();
    t.join();
}
```

One-way communication:  
The thread set the promise and get  
the future

Define the promise and the  
future from the promise

ref generates an object of  
type promise<int> to  
hold a reference to p

# Example

One async task.  
Two way sync

```
#include <future>
using namespace std;
int factorial (std::future<int>& f ) {
    int res = 1;
    int N = f.get();
    for ( int i=N; i> 1; i-- )
        res *=i;
    return res;
}
int main () {
    std::promise<int> p;
    std::future<int> f = p.get_future();
    std::future<int> fu =
        async(std::launch::async, factorial, std::ref(f));
    p.set_value(4);
    int x = fu.get();
}
```

Two-ways communication:  
The caller set the promise and get  
the future

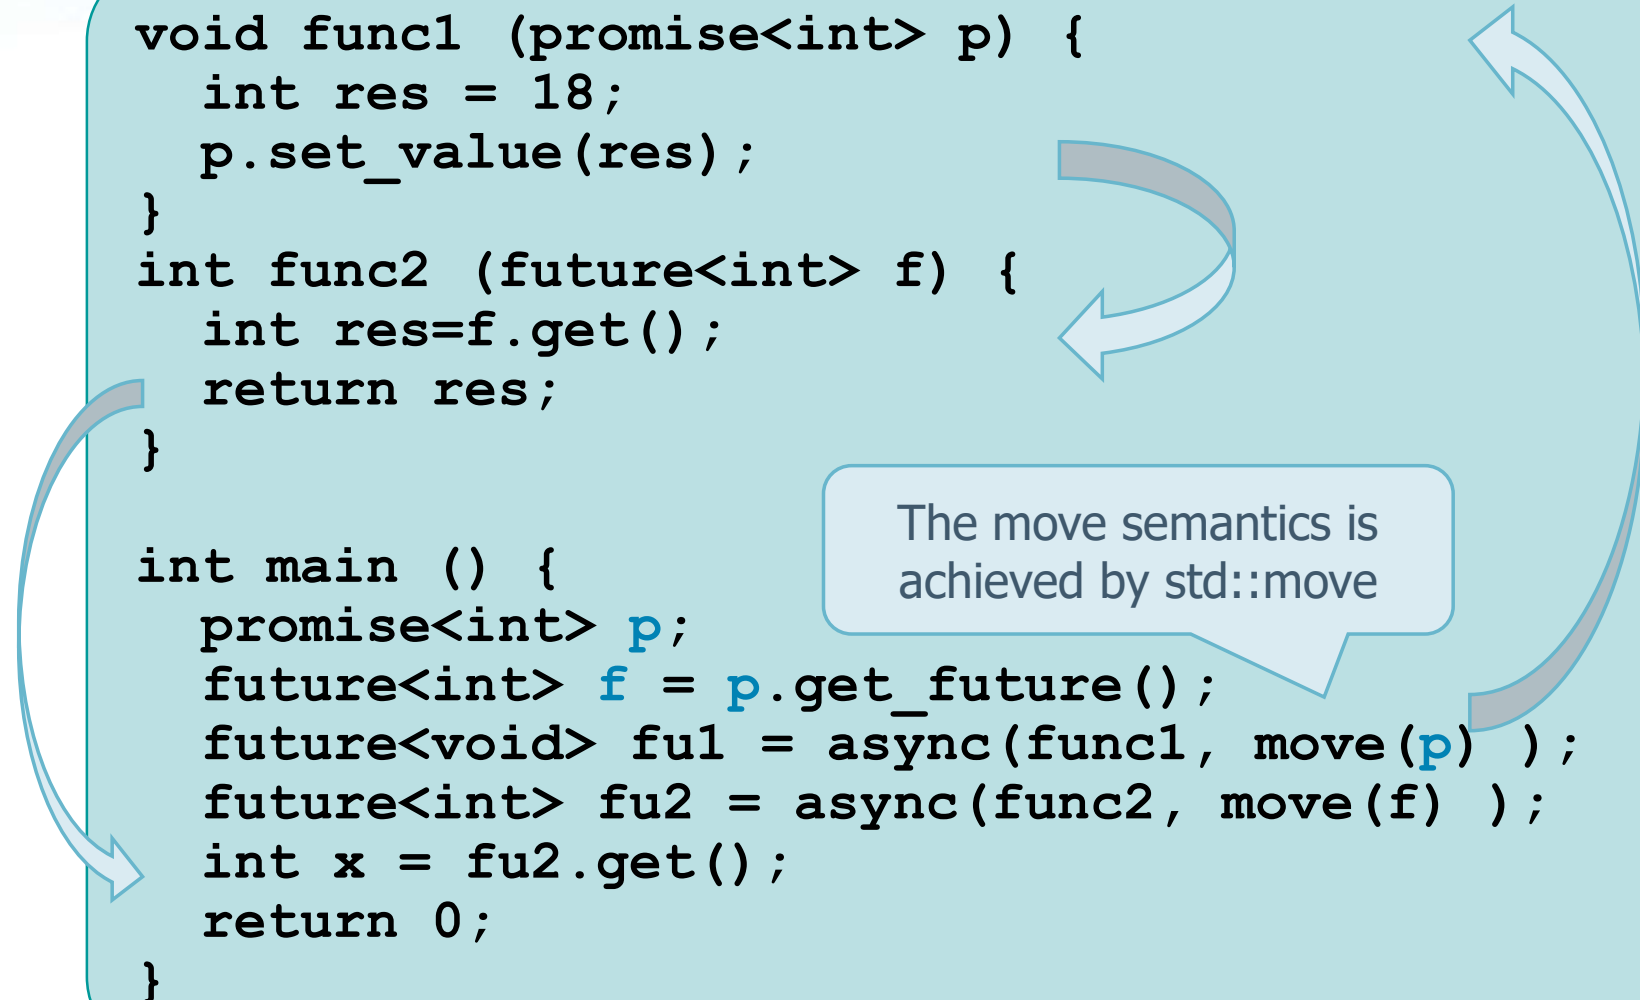
The future must be passed by  
reference, since it doesn't  
support copy semantics

Define the promise and the  
future from the promise

# Example

Two async tasks.  
Two way sync

```
void func1 (promise<int> p) {  
    int res = 18;  
    p.set_value(res);  
}  
int func2 (future<int> f) {  
    int res=f.get();  
    return res;  
}  
  
int main () {  
    promise<int> p;  
    future<int> f = p.get_future();  
    future<void> fu1 = async(func1, move(p) );  
    future<int> fu2 = async(func2, move(f) );  
    int x = fu2.get();  
    return 0;  
}
```



The move semantics is achieved by `std::move`

# Example

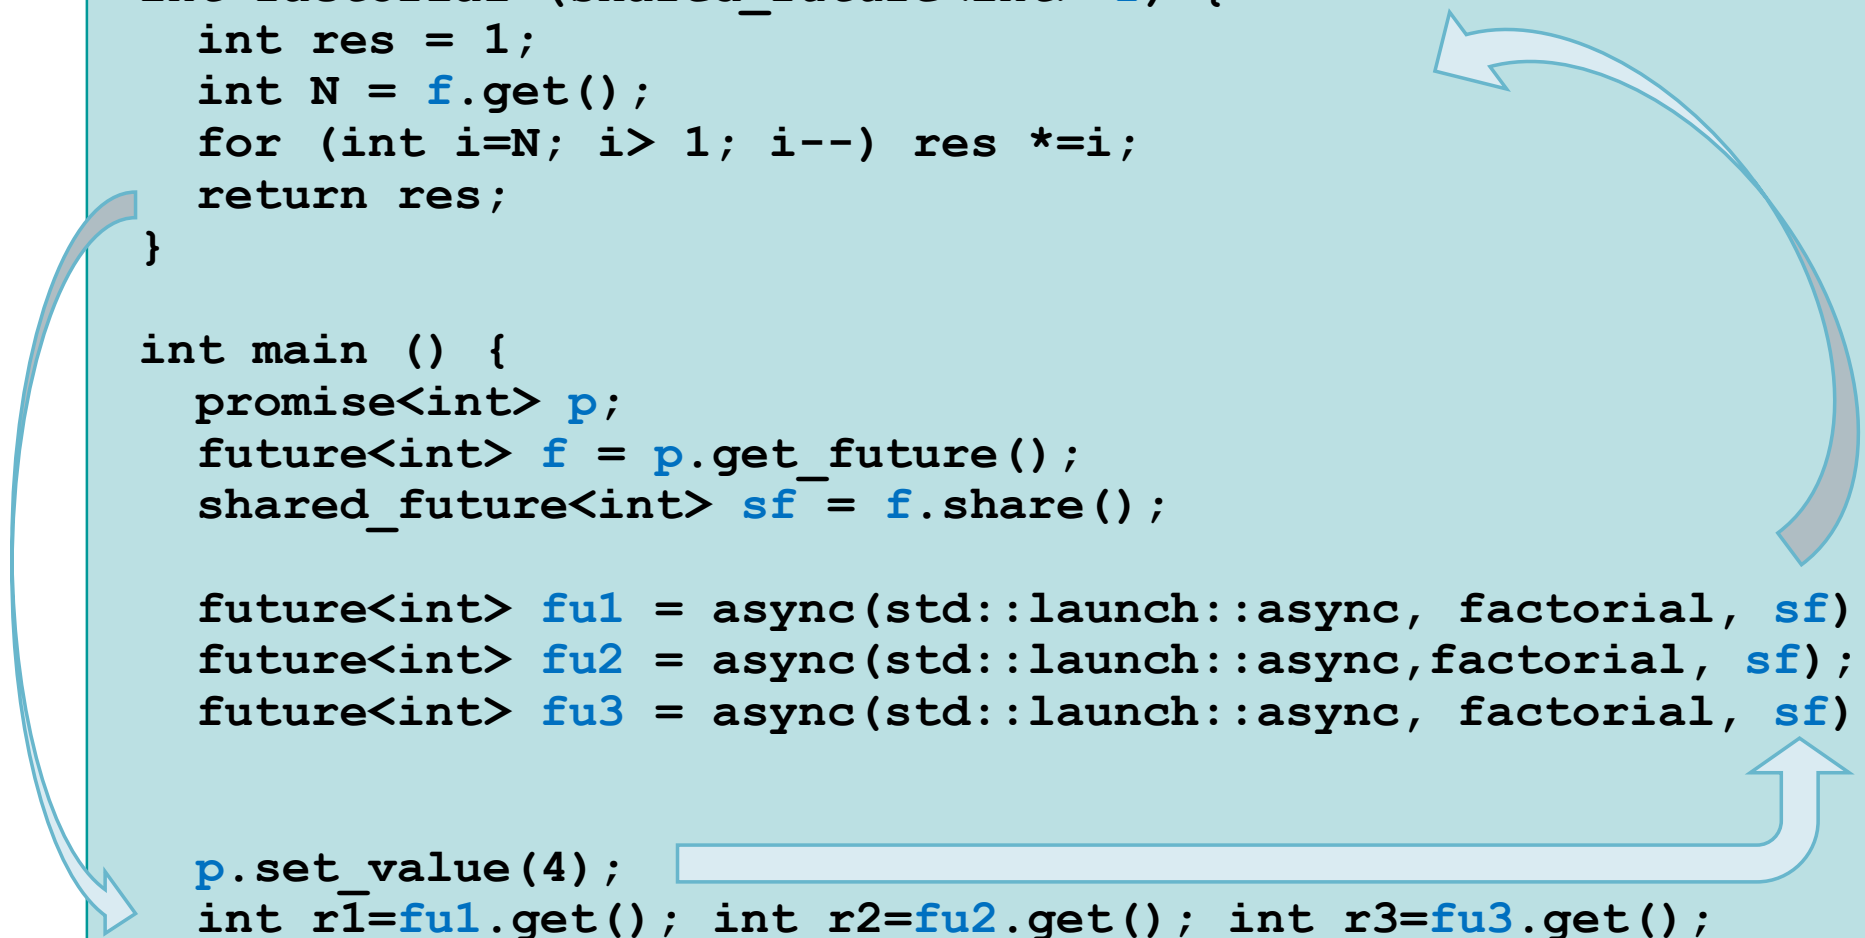
## Shared future

```
using namespace std;
int factorial (shared_future<int> f) {
    int res = 1;
    int N = f.get();
    for (int i=N; i> 1; i--) res *=i;
    return res;
}

int main () {
    promise<int> p;
    future<int> f = p.get_future();
    shared_future<int> sf = f.share();

    future<int> fu1 = async(std::launch::async, factorial, sf);
    future<int> fu2 = async(std::launch::async, factorial, sf);
    future<int> fu3 = async(std::launch::async, factorial, sf);

    p.set_value(4);
    int r1=fu1.get(); int r2=fu2.get(); int r3=fu3.get();
    return 0;
}
```





## Conclusions

### ❖ The task-based approach

- Makes the OS in charge of the parallelism
- Makes the return value of a thread/task accessible
- Run threads with a smart policy
  - CPU load balancing
    - The C++ library can run the function without spawning a thread
  - Avoid the raising of **std::system\_error** in case of thread number reached the system limit
- Allows futures to catch exceptions thrown by the function
  - With **std::thread()** the program terminates

## Conclusions

### ❖ Thread-based approach

- Is used to execute tasks that do not terminate till the end of the application
  - A thread entry point function is like a second, concurrent **main**
- It is a more general concurrency model
  - Can be used for thread-based design patterns
- Allows us to access to the pthread native handle
  - Makes the programmer in charge of the parallelism
  - Useful for advanced management (priority, affinity, scheduling policies, etc.)

## Exercise

- ❖ Resorting only to asynchronous tasks, write a C++ program to perform the matrix multiplication
  - $C = A \times B$
  
- ❖ Constraints
  - The number of columns in *A* must equal the number of rows in *B*
  - Use C++ containers
    - Implement the matrices as vector of vectors
      - `std::vector<std::vector<int>> a, b;`
  - To compute a sum of products, it is possible to use the function **`std::inner_product`**

# Solution

```
#include ...
```

Utility function:  
Generate a random matrix

```
void generateRandomMatrix (  
    vector<vector<int>>& m, int nRow, int nCol){  
    for (int i = 0; i < nRow; i++){  
        vector<int> row;  
        for (int j = 0; j < nCol; j++){  
            row.push_back(rand()%3);  
        }  
        m.push_back(row);  
    }  
}
```

Utility function:  
Display a matrix

```
void printMatrix(const std::vector<std::vector<int>>& m){  
    for (auto & row : m){  
        for (int element : row)  
            cout << element << " ";  
        cout << endl;  
    }  
}
```

# Solution

Function threads

```
int computeSumOfProducts (  
    const std::vector<int>& v1, const std::vector<int>& v2) {  
  
    return std::inner_product (  
        v1.begin(), v1.end(), v2.begin(), 0  
    );  
}
```

Return result

Until C++11

Computes the sum of products on the range [v1.begin, v1.end] and the range beginning at v2.begin, initializing the accumulator at 0

# Solution

Main thread

```
int main() {  
    int nRowA = 2, nColA = 3, nRowB = 3, nColB = 2;  
    vector<std::vector<int>> a, b;  
    vector<std::vector<std::future<int>>> futures;  
    ofstream outputFile;  
  
    cout << "Insert size of matrix A" << endl;  
    cin >> nRowA >> nColA;  
    cout << "Insert size of matrix B" << endl;  
    cin >> nRowB >> nColB;  
  
    if(nColA != nRowB) {  
        cout << "Wrong size colA != rowB" << endl;  
        return -1;  
    }  
}
```

Matrices

Matrix of  
futures

Read the size  
of A and B



# Solution

```
generateRandomMatrix(a, nRowA, nColA);  
// generate the transpose of B (columns -> rows)  
// so that we can easily access the columns of B  
// for the multiplication  
generateRandomMatrix(b, nColB, nRowB);  
  
// Generate the futures to compute the products  
for (int i = 0; i < nRowA; i++){  
    vector<std::future<int>> futureRow;  
    for (int j = 0; j < nColB; j++) {  
        future<int> f = std::async(  
            std::launch::async | std::launch::deferred,  
            computeSumOfProducts, a[i], b[j]);  
        // Futures are not copyable so we have to use move  
        // to store them in a vector  
        futureRow.push_back(std::move(f));  
    }  
    futures.push_back(std::move(futureRow));  
}
```

Generate matrix A

Generate matrix B

Create a future for  
each element

Insert futures in the  
matrix of futures

# Solution

```
cout << "A" << endl;  
printMatrix(a);  
cout << "B" << std::endl;  
printMatrix(b);
```

Display input matrices

```
outputFile.open("./output-matrix.txt");  
for (auto & row : futures){  
    for (std::future<int> & f : row) {  
        outputFile << f.get() << " ";  
    }  
    outputFile << std::endl;  
}  
outputFile.close();  
  
return 0;  
}
```

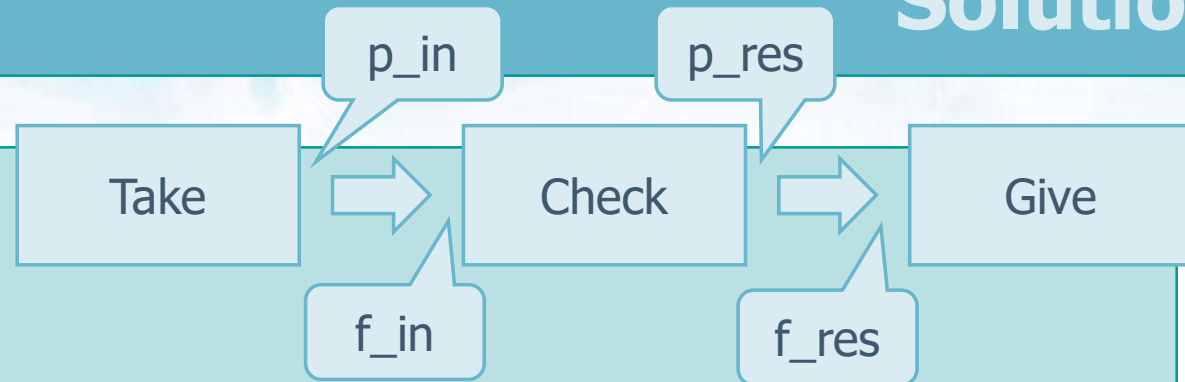
Wait futures to be ready  
and display result

## Exercise

Exam 5 July 2021

- ❖ Write a C++ program with three tasks
  - Thread **take** reads a number from command line
  - Thread **check** checks whether the number is prime
  - Thread **give** displays the answer to standard output
- ❖ Thread communication should be made using promises and futures
  - All functions are acyclic

# Solution

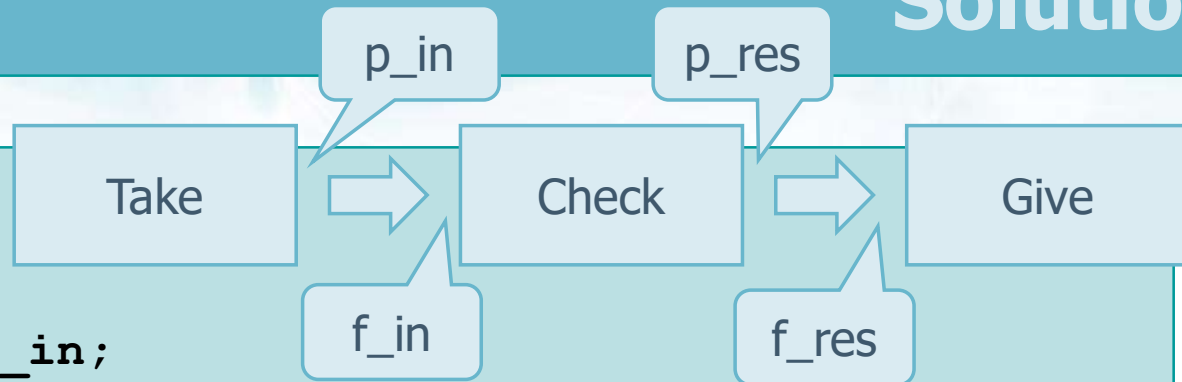


```
#include <iostream>
#include <thread>
#include <vector>
#include <future>
```

```
void take (std::promise<int>&);
void check (std::future<int>&, std::promise<bool>&);
void give (std::future<bool>&);
```

Thread functions

# Solution



```
int main() {  
    std::promise<int> p_in;  
    std::future<int> f_in = p_in.get_future();  
  
    std::promise<bool> p_res;  
    std::future<bool> f_res = p_res.get_future();  
  
    std::thread t1(take, std::ref(p_in));  
    std::thread t2(check, std::ref(f_in), std::ref(p_res));  
    std::thread t3(give, std::ref(f_res));  
  
    t1.join();  
    t2.join();  
    t3.join();  
    return 0;  
}
```

# Solution

Reading thread

```
void take (std::promise<int> &p_in) {  
    int in;  
    std::cout << "Insert a number" << std::endl;  
    std::cin >> in;  
    p_in.set_value (in);  
}
```

Set promise "in"

Writing thread

```
void give (std::future<bool>& f_res) {  
    bool answer = f_res.get();  
    std::string s0 (" ");  
    if(!answer)  
        s0=" NOT";  
    std::cout << "Number is" << s0 << " prime";  
}
```

Get future "ref"

# Solution

Computation thread

```
void check (  
    std::future<int> &f_in, std::promise<bool>& p_res)  
{  
    int n = f_in.get();  
    bool prime=true;  
    if (n <= 1){  
        prime = false;  
    }  
    // Check from 2 to n-1  
    for (int j=2; j<n; j++) {  
        if (n % j == 0) {  
            prime = false;  
            break;  
        }  
    }  
    p_res.set_value(prime);  
}
```

Get future "in"

Set promise  
"res"