

OS161

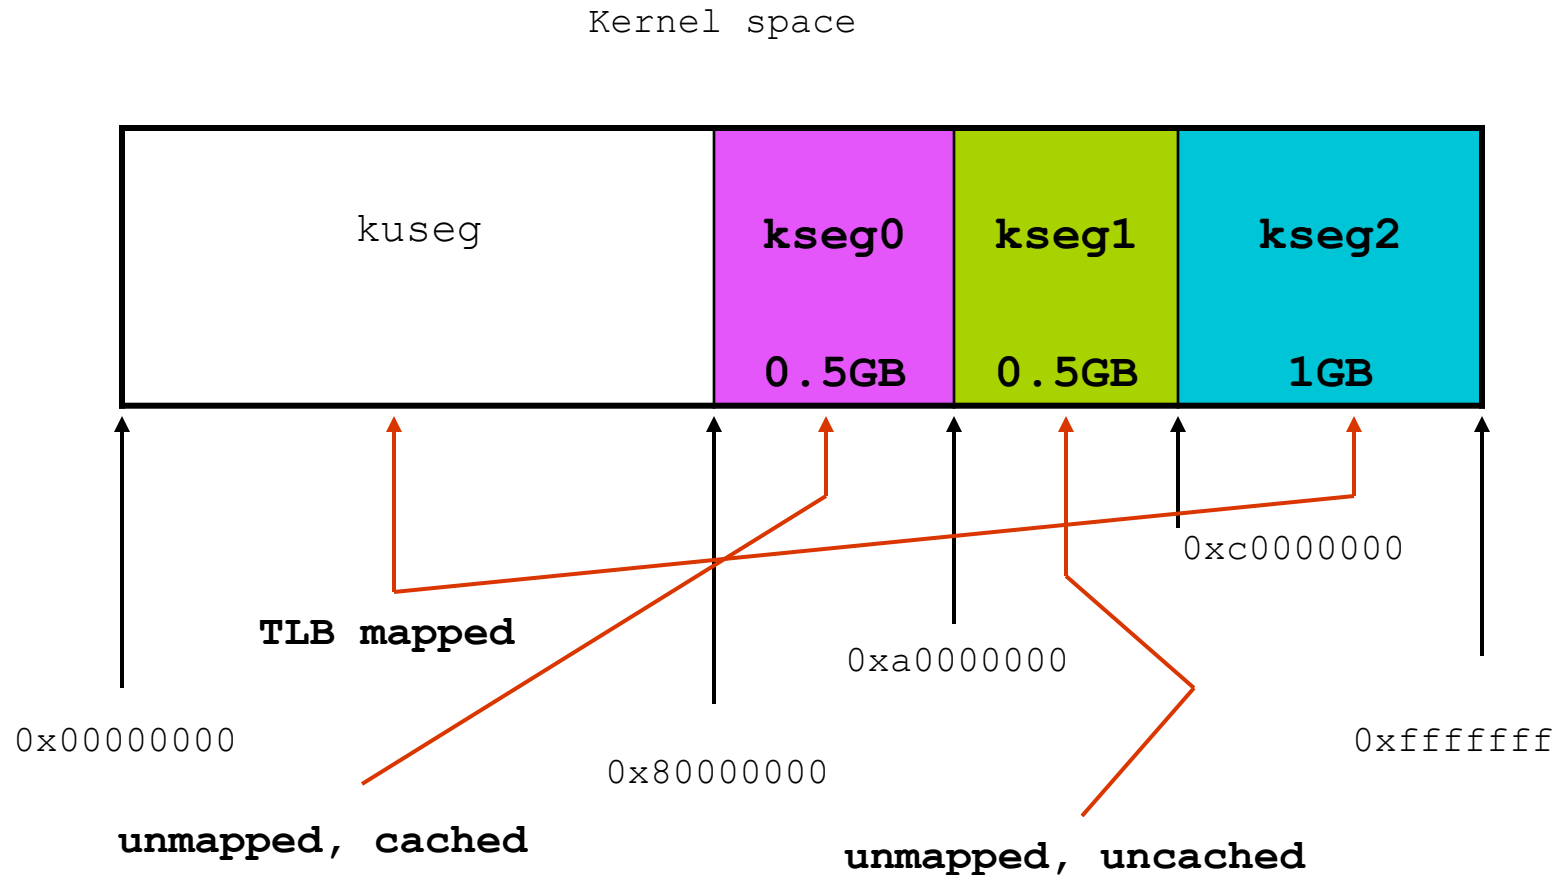
Address Space & Memory Management

Dumbvm and kmalloc

- Contiguous allocation
- Allocating by page multiples (4096 byte frame)
 - getppages (dumbvm.c): calls ram_stealmem (in mutual exclusion)
 - ram_stealmem (ram.c): allocates contiguous RAM starting at firstpaddr, that is increased
- Allocator is common to
 - User memory: as_prepare_load calls getppages for 2 user segments and a stack
 - Dynamic kernel memory: kmalloc is based on alloc_kpages, that calls getppages

mips is a 32 bits microprocessor means 4gigs, 2 for kernel and 2 for? Difference is the first bit, the most significant bit in the 32 bits. So kuseg is the user segment which is user memory. There are microprocessors 1000 for users and 1000 for processes.

MIPS VIRTUAL ADDRESS SPACE



In OS/161, user programs live in kuseg, kernel code and data structures live in kseg0, devices are accessed through kseg1, and kseg2 is not used.

Kernel loader (sys161: start.S)

Logical addr. (KSEG0)

Physical addr.

0x80000000

0x0

exception handlers

0x80000200

0x200

kernel

0x80039d54 (`_end`)

0x39d54

arg string for boot +
Page align

0x8003a000 (P)

0x3a000

Stack for first thread
(1 page = 4096 B)

0x8003b000 (P+1000)

0x3b000

FREE MEMORY

0x80100000

0x100000

ramsize (es. 1MB: sys161.conf)

Kernel loader (sys161: start.S)

Logical addr. (KSEG0)

Physical addr.

0x80000000

0x0

exception handlers

0x80000200

0x200

kernel

0x80039d54 (`_end`)

0x39d54

arg string for boot +
Page align

0x8003a000 (P)

0x3a000

Stack for first thread
(1 page = 4096 B)

0x8003b000

0x3b000

FREE MEMORY

(firstfree)

(firstpaddr)

0x80100000

ramsize (es. 1MB: sys161.conf)

0x100000

Dumbvm

Logical addr. (KSEG0)

Physical addr.

0x80000000

0x0

0x8003b000

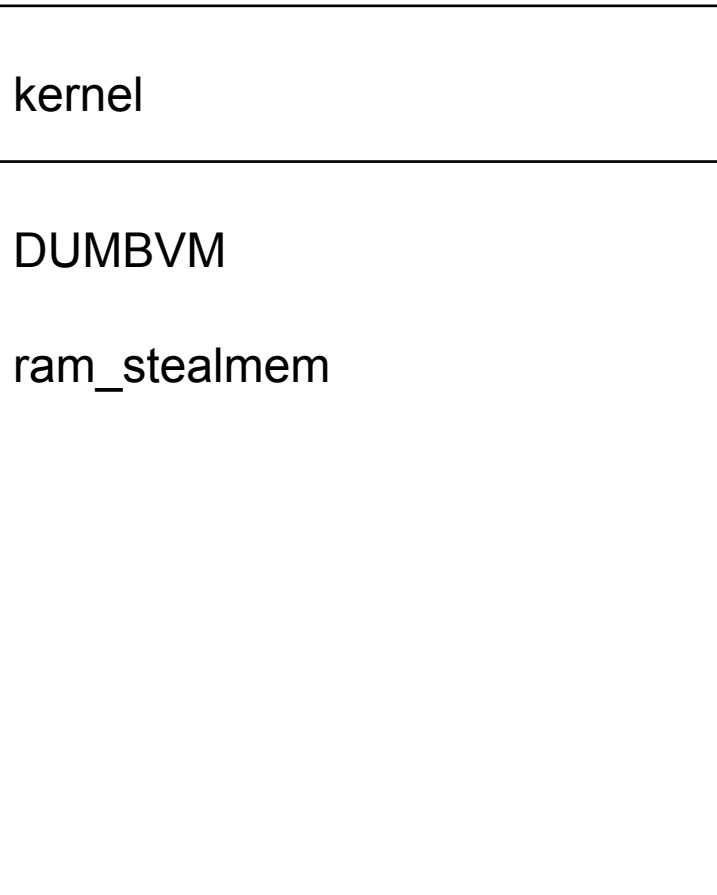
(firstfree)

0x3b000

(firstpaddr)

0x80100000

0x100000



ramsize

ram_bootstrap

```
void
ram_bootstrap(void) {
    /* Get size of RAM. */
    size_t ramsize = mainbus_ramsize();
    if (ramsize > 512*1024*1024) {
        ramsize = 512*1024*1024;
    }
    lastpaddr = ramsize;
    /* Get first free virtual address from where
       start.S saved it. Convert to physical address. */
    firstpaddr = firstfree - MIPS_KSEG0;
}
```

ram_stealmem

(kern/arch/mips/vm/ram.c)

```
paddr_t ram_stealmem(unsigned long npages) {  
    paddr_t paddr;  
    size_t size = npages * PAGE_SIZE;  
  
    if (firstpaddr + size > lastpaddr) {  
        return 0;  
    }  
  
    paddr = firstpaddr;  
    firstpaddr += size;  
  
    return paddr;  
}
```


getppages

(kern/arch/mips/vm/dumbvm.c)

```
static paddr_t
getppages(unsigned long npages) {
    paddr_t addr;

    spinlock_acquire(&stealmem_lock);

    addr = ram_stealmem(npages);

    spinlock_release(&stealmem_lock);

    return addr;
}
```

getppages

(kern/arch/mips/vm/dumbvm.c)

```
static paddr_t  
getppages(unsigned long npages) {  
    paddr_t addr;  
  
    spinlock_acquire(&stealmem_lock);  
  
    addr = ram_stealmem(npages);  
  
    spinlock_release(&stealmem_lock);  
  
    return addr;  
}
```

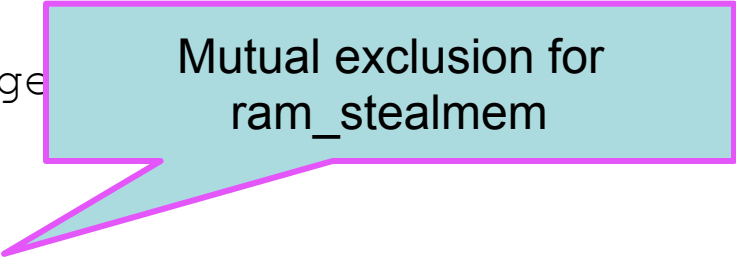


Internal (dumbvm) function

getppages

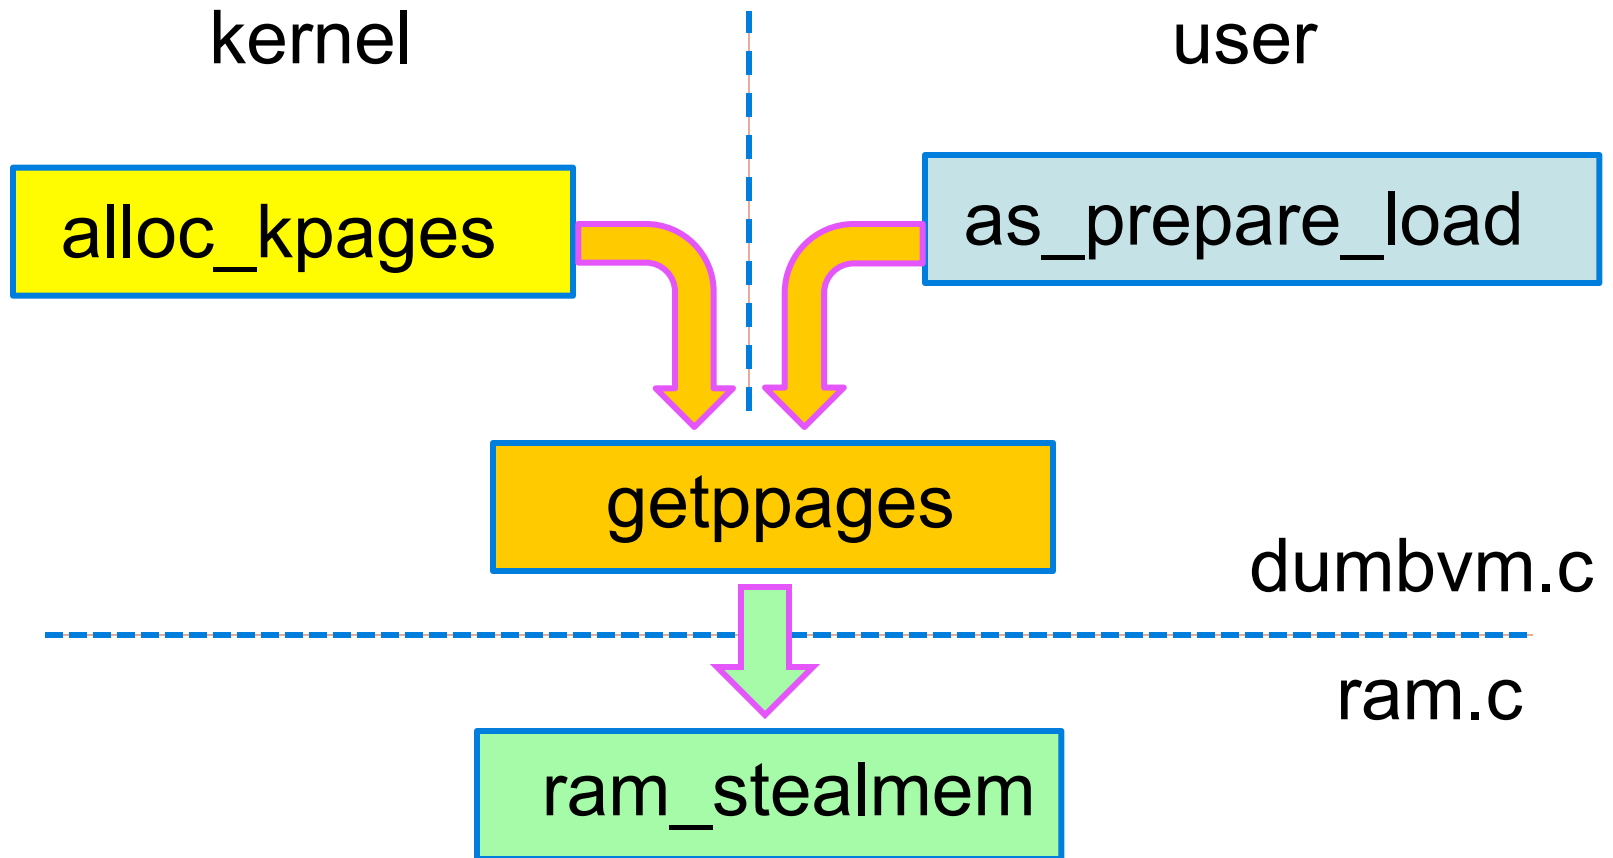
(kern/arch/mips/vm/dumbvm.c)

```
static paddr_t  
getppages(unsigned long npages,  
           paddr_t addr;  
  
           spinlock_acquire(&stealmem_lock);  
  
           addr = ram_stealmem(npages);  
  
           spinlock_release(&stealmem_lock);  
  
           return addr;  
}
```



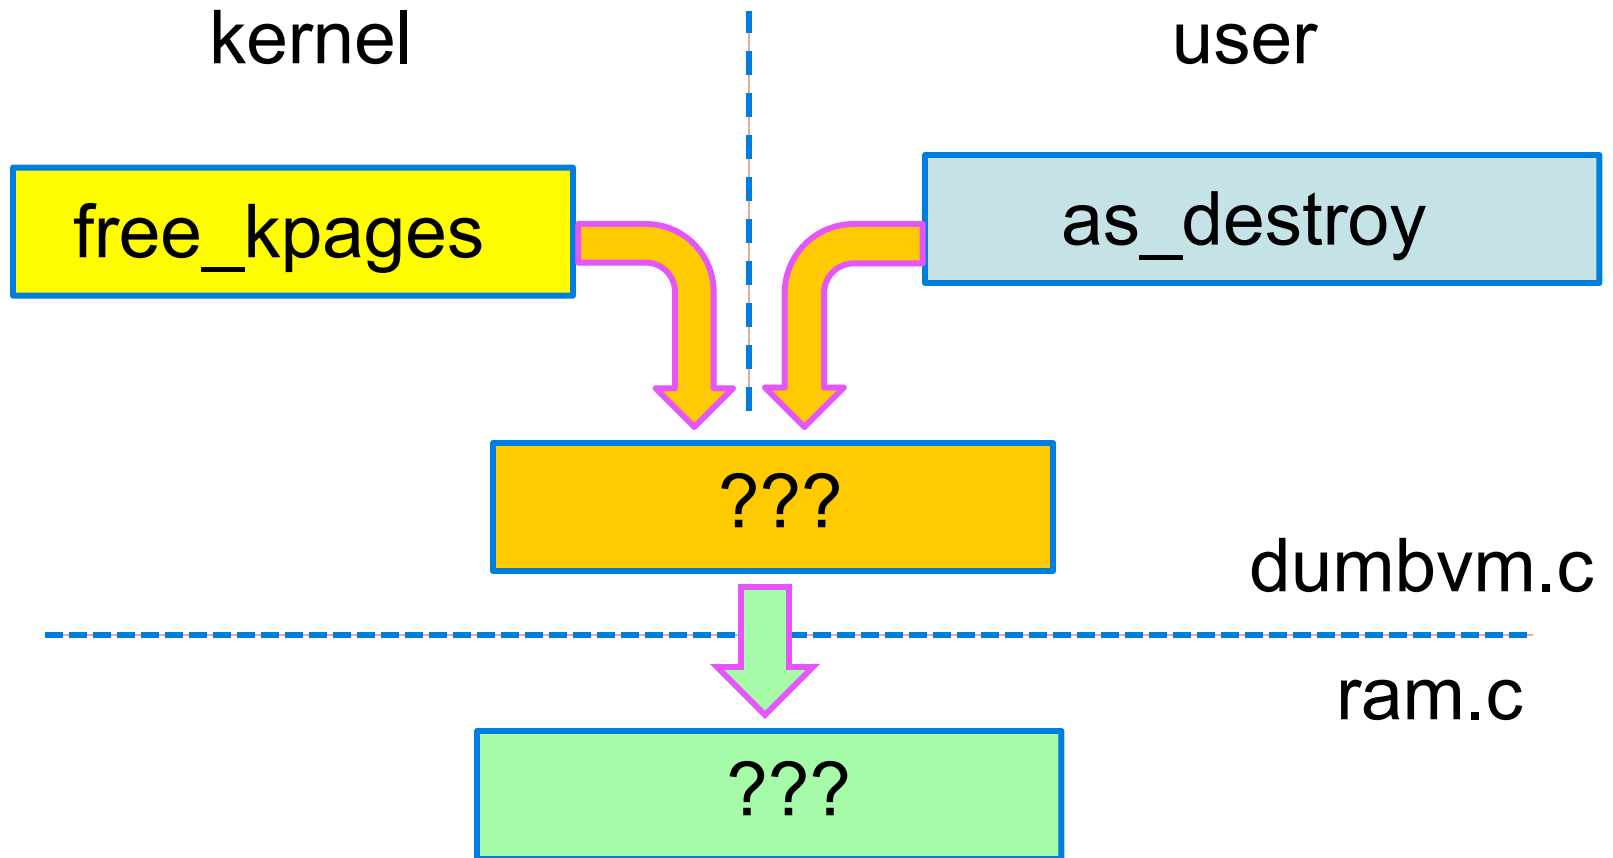
Mutual exclusion for
ram_stealmem

dumbvm.c (alloc)



dumbvm.c (free)

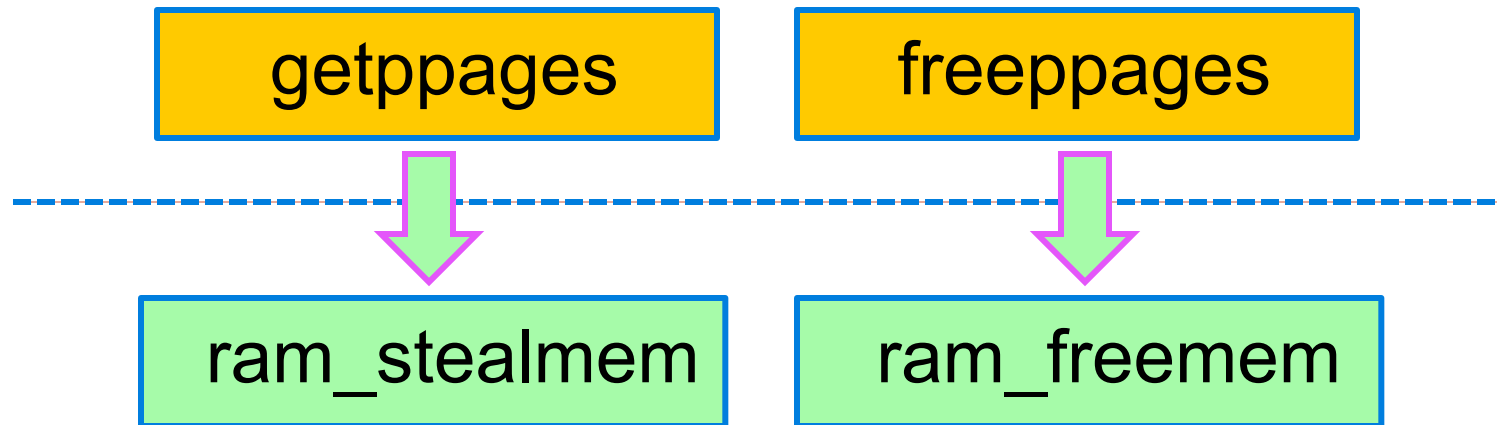
not implemented -> TO DO!



De-alloc (free) in ram.c

(solution A1)

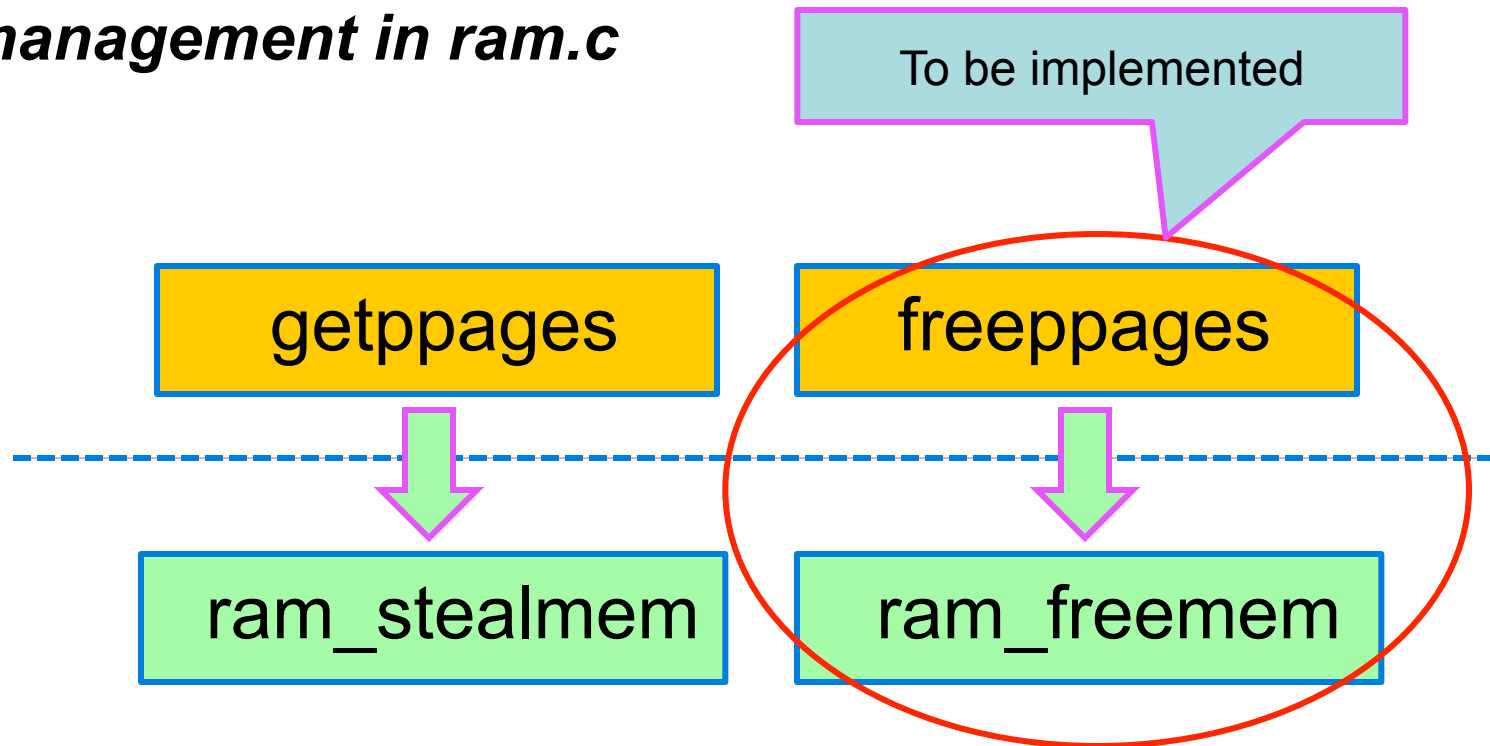
- freeppages just an interface to ram_freemem
- Data structure (free-list or bitmap) and **memory management in ram.c**



De-alloc (free) in ram.c

(*solution A1*)

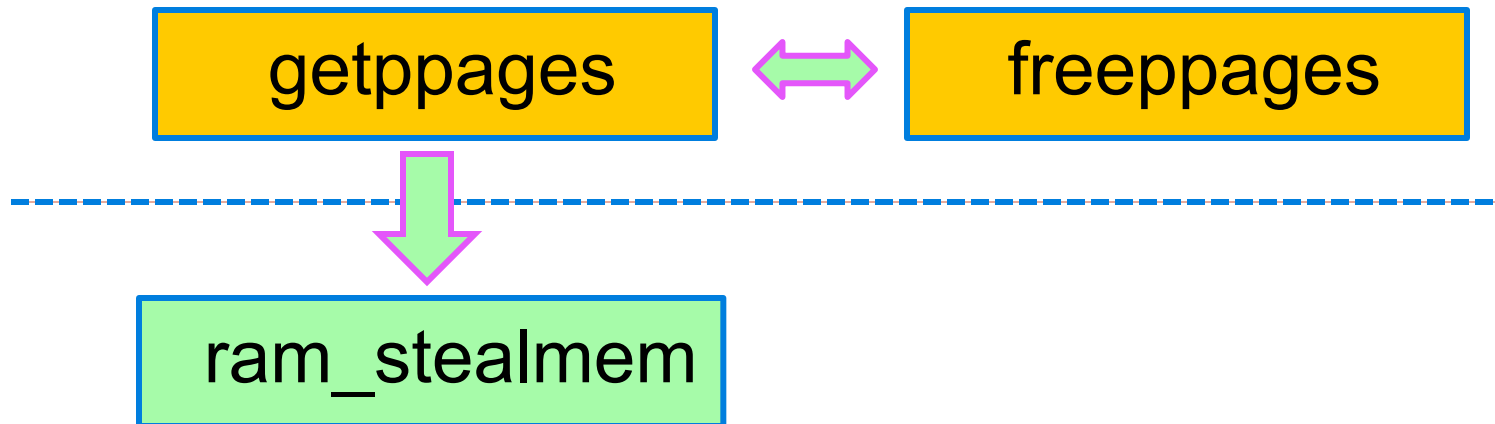
- freeppages just an interface to ram_freemem
- Data structure (free-list or bitmap) and **memory management in ram.c**



De-alloc (free) in ram.c

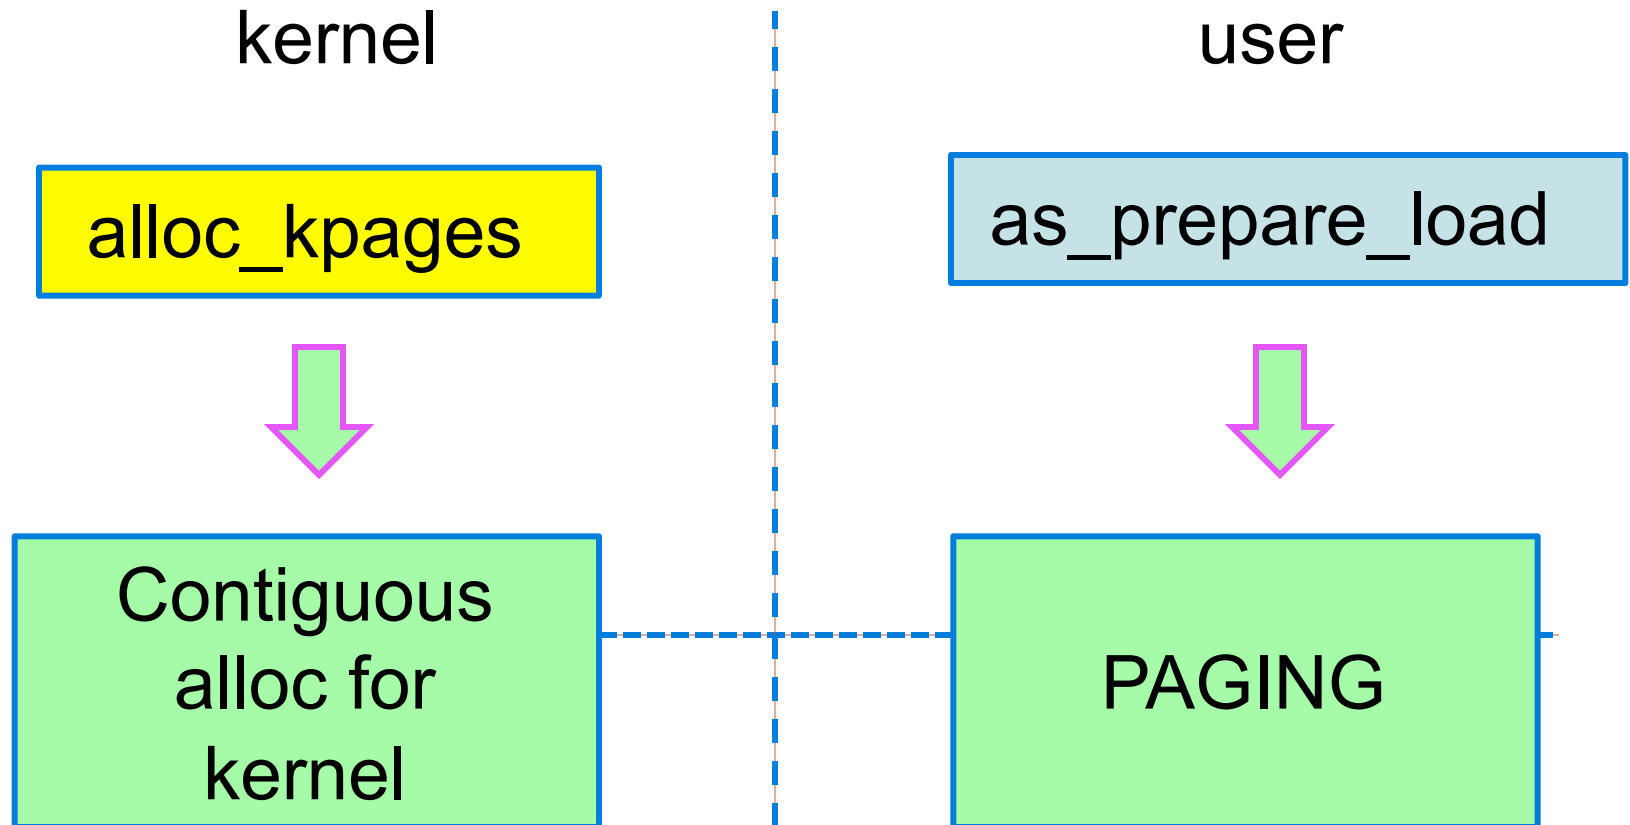
(*solution A2*)

- Memory not returned to RAM
- Data structure (free-list or bitmap) and **memory management in dumbvm.c**
- Freeppages coordinates with getppages



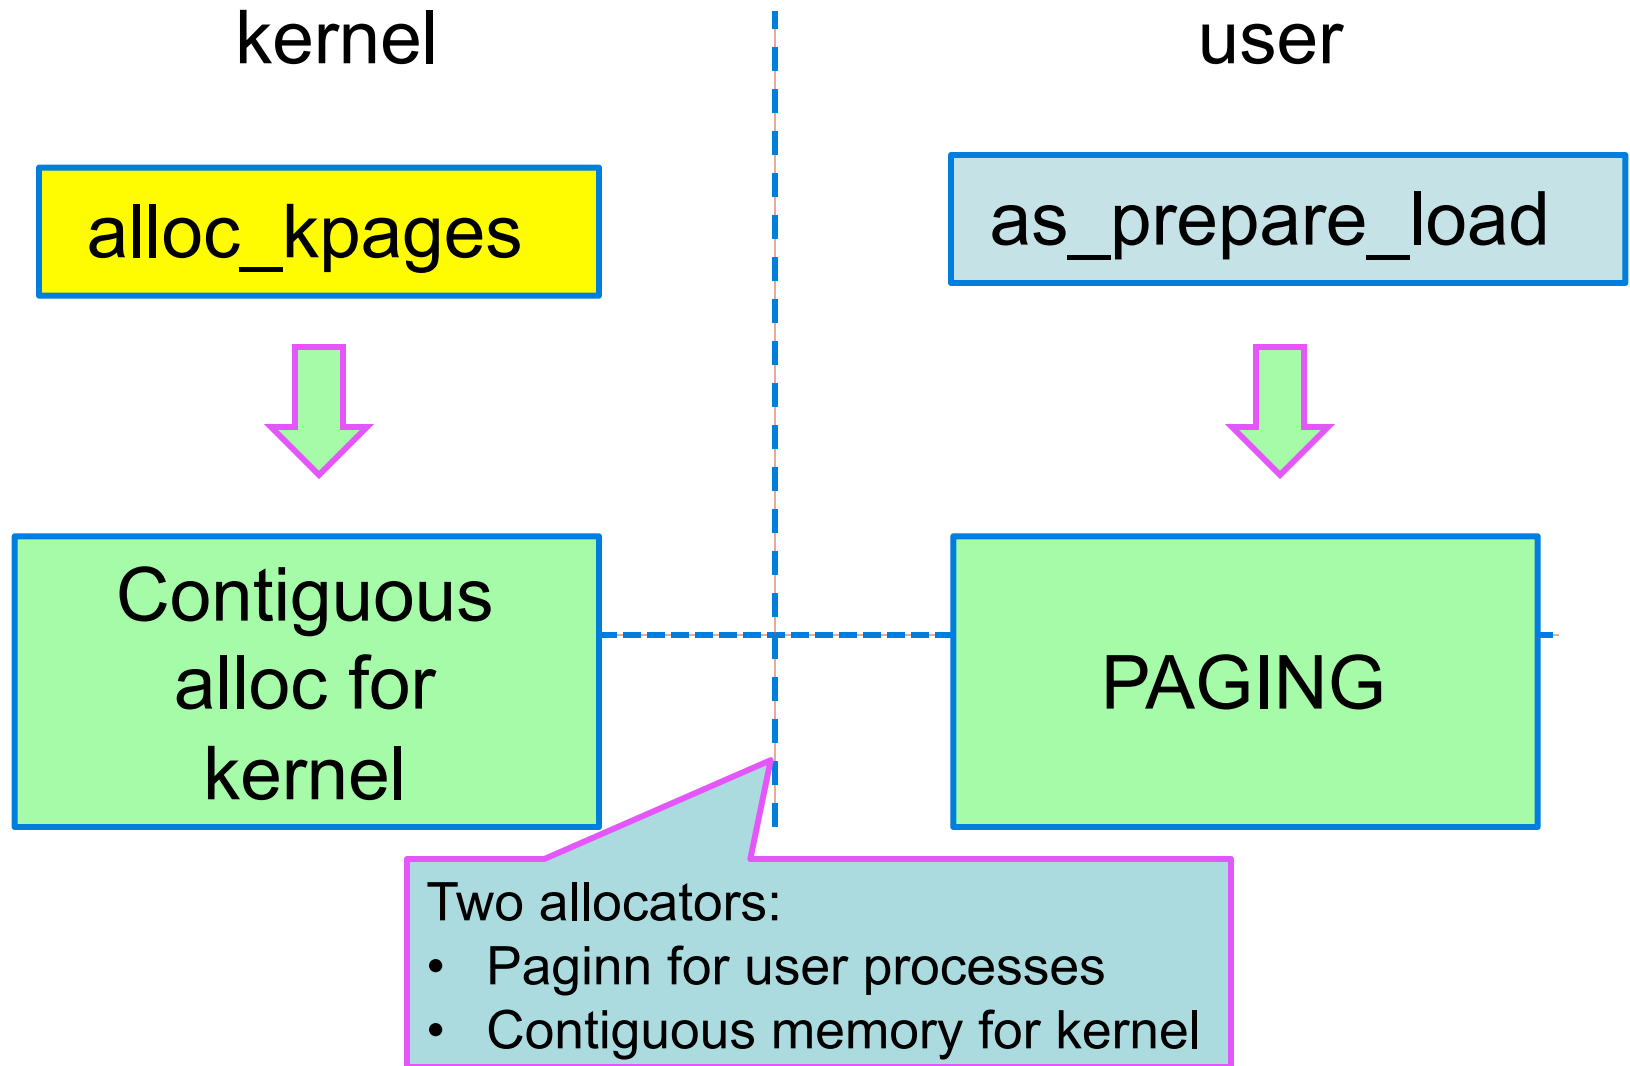
Paging in user space

(solution B)



Paging in user space

(solution B)



Proposed solution

de-alloc in dumbvm (sol. A2)

- Contiguous allocation(by pages) common to kernel and user
- Allocator in dumbvm: keep track (using a bitmap) of previously freed pages. In order to alloc
 - First search among (previously) freed pages (**an interval of contiguous free pages**)
 - **If not found, call ram_stealmem**
- Bitmap implemented as an array of char (for simplicity)
 - `freeRamFrames[i] = 1/0` (free/alloated): free=**FREED!** (by **freepages**)
- In order to free we need to know
 - Pointer (or index) to first page in interval
 - Size, i.e. number of (contiguous) pages to free
- We need a table to store sizes (number of pages in allocated intervals) for each alloc performed
 - `void free_kpages(vaddr_t addr)`: table needed as only pointer passed
 - `void as_destroy(struct addrspace *as)`: table not needed as size is stored in address space
- `allocSize[i] = /* number of pages allocated starting at i-th */`

Global variables (and test function)

```
static struct spinlock freemem_lock = SPINLOCK_INITIALIZER;

static unsigned char *freeRamFrames = NULL;
static unsigned long *allocSize = NULL;
static int nRamFrames = 0;

static int allocTableActive = 0;

static int isTableActive () {
    int active;
    spinlock_acquire(&freemem_lock);
    active = allocTableActive;
    spinlock_release(&freemem_lock);
    return active;
}
```

Global variables (and test function)

```
static struct spinlock freemem_lock = SPINLOCK_INITIALIZER;

static unsigned char *freeRamFrames = NULL;
static unsigned long *allocSize = NULL;
static int nRamFrames = 0;

static int allocTableSize = 0;

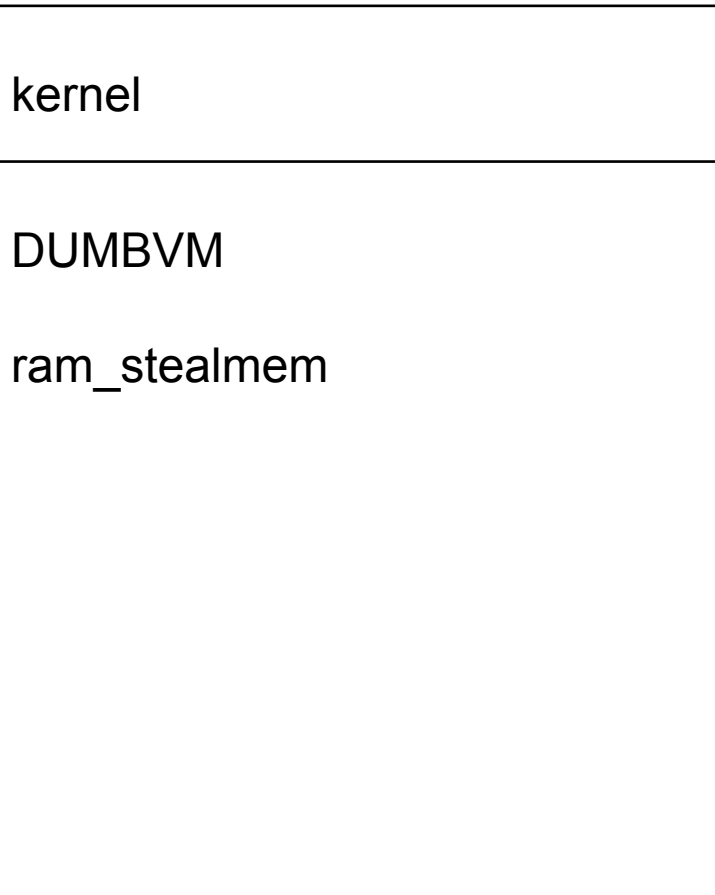
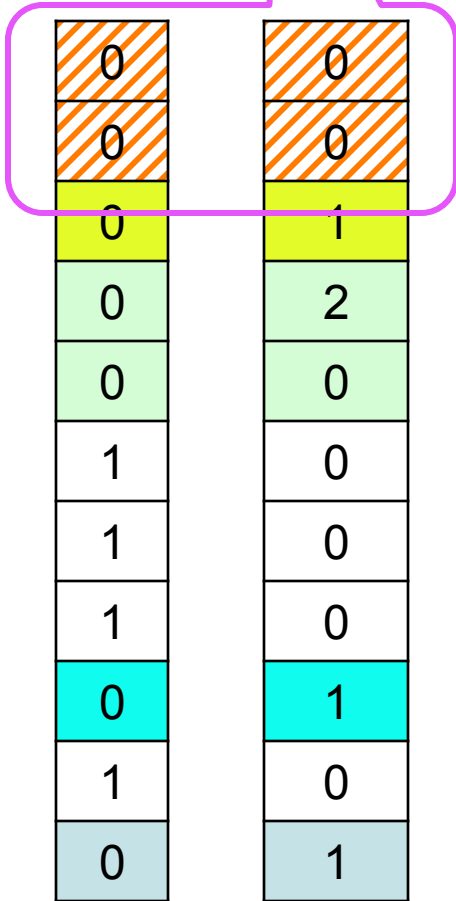
static int
    int active;
    spinlock_t lock;
    active = 0;
    spinlock_release(&freemem_lock);
    return active;
}
```

Dynamic arrays as RAM size known
at Boot (depends on da sys161.conf)
Alternative: over-dimensioned static arrays!

Dumbvm

Never freed

freeRamFrames



Physical addr.

0x0

0x3b000

0x100000

Never freed:
including allocations
before table available

Dumbvm

freeRamFrames

Physical addr.

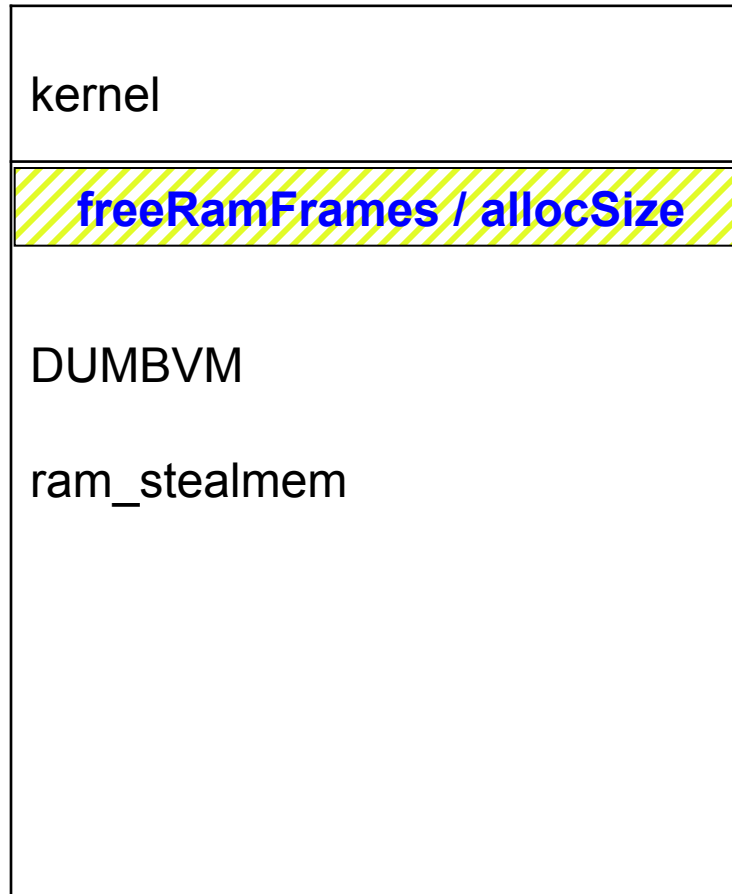
0x0

0x3b000

0x100000

0	0
0	0
0	0
0	2
0	0
1	0
1	0
1	0
0	1
1	0
0	1

allocSize



Dumbvm

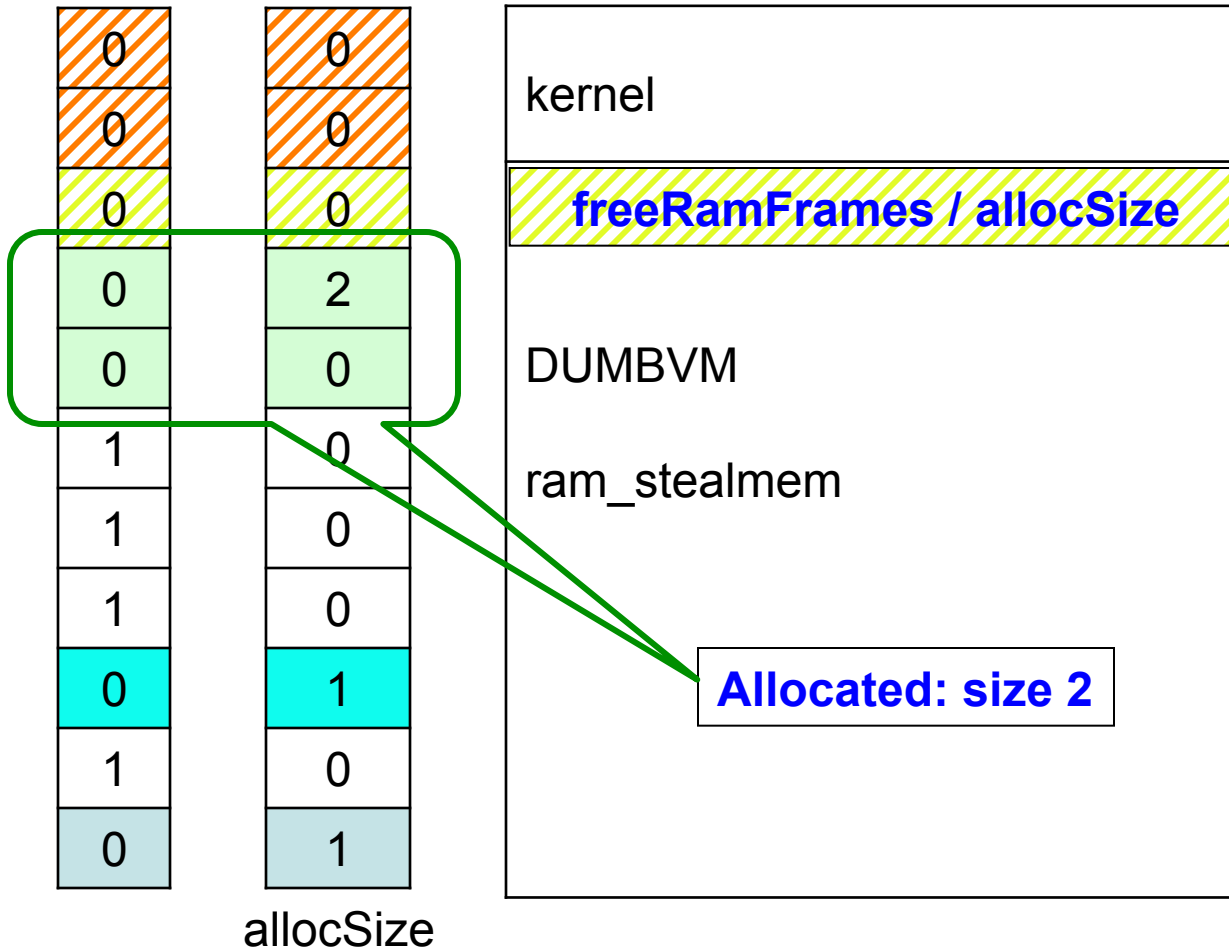
freeRamFrames

Physical addr.

0x0

0x3b000

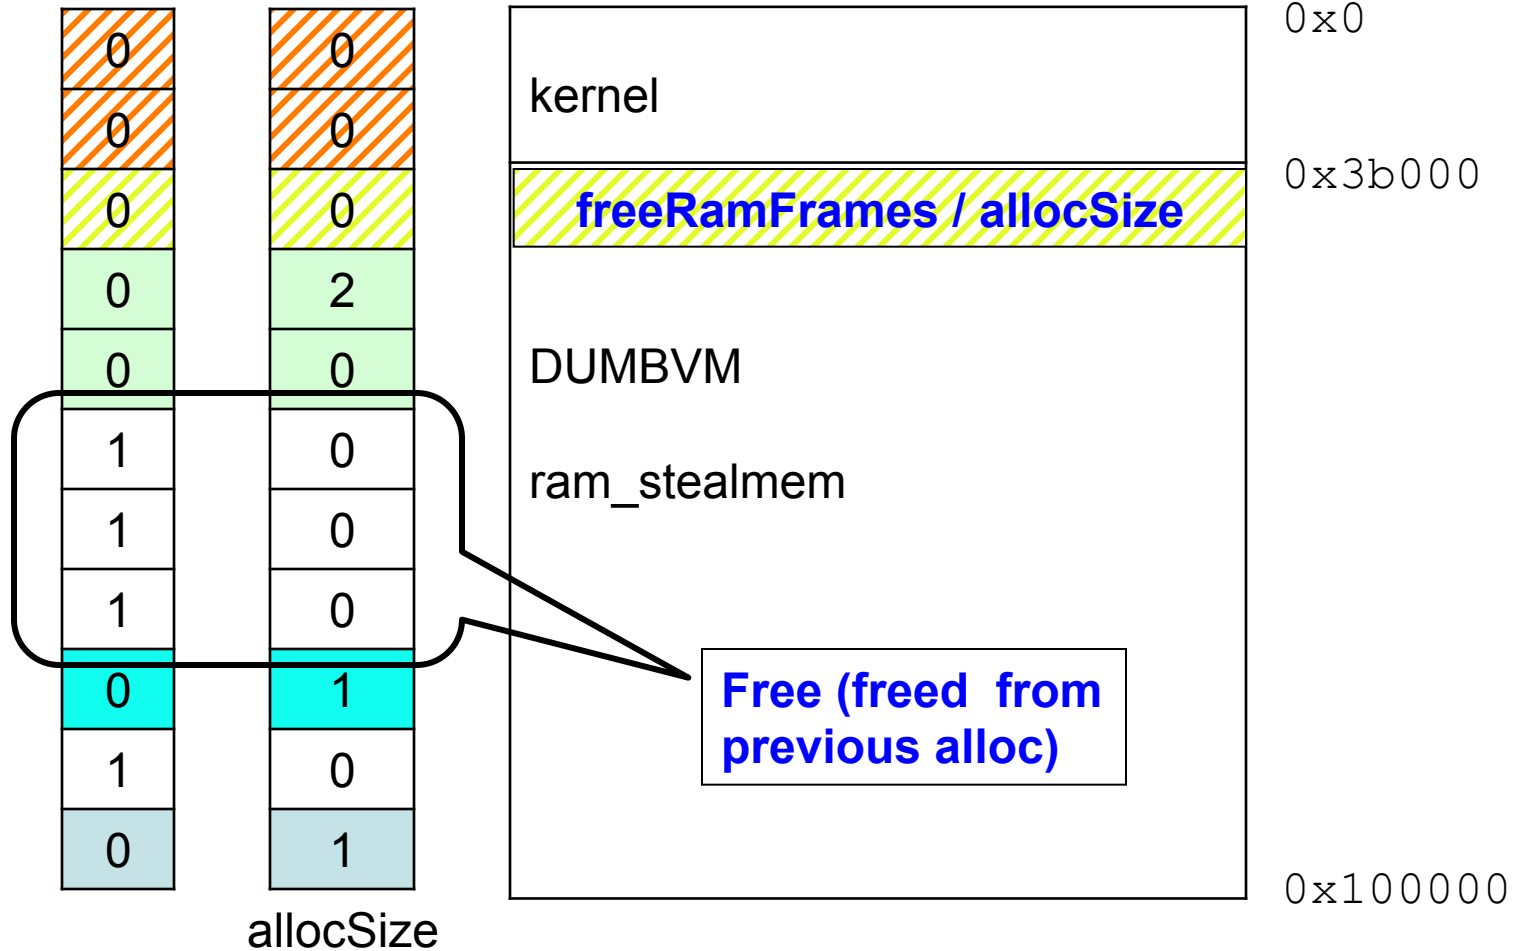
0x100000



Dumbvm

freeRamFrames

Physical addr.



free_kpages & as_destroy

```
void free_kpages(vaddr_t addr){
    if (isTableActive()) {
        paddr_t paddr = addr - MIPS_KSEG0;
        long first = paddr/PAGE_SIZE;
        KASSERT(nRamFrames>first);
        freeppages(paddr, allocSize[first]);
    }
}

void as_destroy(struct addrspace *as){
    dumbvm_can_sleep();
    freeppages(as->as_pbase1, as->as_npages1);
    freeppages(as->as_pbase2, as->as_npages2);
    freeppages(as->as_stackbase, DUMBVM_STACKPAGES);
    kfree(as);
}
```

Initialization

```
void vm_bootstrap(void) {
    int i;
    nRamFrames = ((int)ram_getsize())/PAGE_SIZE;
    /* alloc freeRamFrame and allocSize */
    freeRamFrames = kmalloc(sizeof(unsigned char)*nRamFrames);
    allocSize      = kmalloc(sizeof(unsigned long)*nRamFrames);
    if (freeRamFrames==NULL || allocSize==NULL) {
        /* reset to disable this vm management */
        freeRamFrames = allocSize = NULL; return;
    }
    for (i=0; i<nRamFrames; i++) {
        freeRamFrames[i] = (unsigned char)0; allocSize[i] = 0;
    }
    spinlock_acquire(&freemem_lock);
    allocTableActive = 1;
    spinlock_release(&freemem_lock);
}
```

getppages

```
static paddr_t getppages(unsigned long npages) {
    paddr_t addr;

    /* try freed pages first */
    addr = getfreeppages(npages);
    if (addr == 0) { /* call stealmem */
        spinlock_acquire(&stealmem_lock);
        addr = ram_stealmem(npages);
        spinlock_release(&stealmem_lock);
    }
    if (addr != 0 && isTableActive()) {
        spinlock_acquire(&freemem_lock);
        allocSize[addr/PAGE_SIZE] = npages;
        spinlock_release(&freemem_lock);
    }
    return addr;
}
```

getfreeppages

```
static paddr_t getfreeppages(unsigned long npages) {
    paddr_t addr;
    long i, first, found, np = (long)npages;

    if (!isTableActive()) return 0;
    spinlock_acquire(&freemem_lock);
    // Linear search of free interval
    for (i=0, first=found=-1; i<nRamFrames; i++) {
        if (freeRamFrames[i]) {
            if (i==0 || !freeRamFrames[i-1])
                first = i; /* set first free in an interval */
            if (i-first+1 >= np)
                found = first;
        }
    }
}
```

getfreeppages

```
if (found>=0) {
    for (i=found; i<found+np; i++) {
        freeRamFrames[i] = (unsigned char)0;
    }
    allocSize[found] = np;
    addr = (paddr_t) found*PAGE_SIZE;
}
else {
    addr = 0;
}

spinlock_release(&freemem_lock);

return addr;
}
```

freeppages

```
static int freeppages(paddr_t addr, unsigned long npages){
    long i, first, np=(long)npages;
    if (!isTableActive()) return 0;
    first = addr/PAGE_SIZE;
    KASSERT(allocSize!=NULL);
    KASSERT(nRamFrames>first);

    spinlock_acquire(&freemem_lock);
    for (i=first; i<first+np; i++) {
        freeRamFrames[i] = (unsigned char)1;
    }
    spinlock_release(&freemem_lock);

    return 1;
}
```